Project Report

on

An Image-Based Study on Crop Disease Detection

submitted as partial fulfillment for the award of

BACHELOR OF TECHNOLOGY DEGREE

SESSION 2025-26

in

Computer Science & Engineering

By

Mohammad Areeb(2200290100103)

Ayush Tripathi(2200290100051)

Aman Singh Bohra(2200290100017)

Yash Vinayak Singh(2200290100197)

Under the supervision of

Mr. Rahul Kumar Sharma

KIET Group of Institutions, Ghaziabad

Affiliated to

Dr. A.P.J. Abdul Kalam Technical University, Lucknow

(Formerly UPTU)

May, 2026

DECLARATION

We hereby declare that this submission is our own work and that, to the best of our knowledge and belief, it contains no material previously published or written by another person nor material which to a substantial extent has been accepted for the award of any other degree or diploma of the university or other institute of higher learning, except where due acknowledgment has been made in the text.

Signature:

Name(s):Mohammad Areeb/Ayush Tripathi/Aman Singh Bohra/Yash Vinayak Singh

Roll No.(s):2200290100103/2200290100051/2200290100017/2200290100197

Date:

CERTIFICATE

This is to certify that Project Report entitled "An Image-Based Study on Crop Disease Detection" which is submitted by "Mohammad Areeb/Ayush Tripathi/Yash Vinayak Singh/Aman Singh Bohra" in partial fulfillment of the requirement for the award of degree B. Tech. in Department of Computer Science & Engineering of Dr. A.P.J. Abdul Kalam Technical University, Lucknow is a record of the candidates own work carried out by them under my supervision. The matter embodied in this report is original and has not been submitted for the award of any other degree.

Mr. Rahul Kumar Sharma	**Dr. Vineet Kumar Sharma**
(Assistant Professor)	**(Dean, CSE)**

Date:

ACKNOWLEDGEMENT

It gives us a great sense of pleasure to present the report of the B. Tech Project undertaken during B. Tech. Final Year. We owe special debt of gratitude to Mr. Rahul Kumar Sharma (Assistant Professor), Department of Computer Science & Engineering, KIET Group of Institutions, Ghaziabad, for his constant support and guidance throughout the course of our work. His sincerity, thoroughness and perseverance have been a constant source of inspiration for us. It is only his cognizant efforts that our endeavors have seen light of the day.

We also take the opportunity to acknowledge the contribution of Dr. Vineet Sharma, Dean, Computer Science & Engineering Department, KIET Group of Institutions, Ghaziabad, for his full support and assistance during the development of the project.

We also do not like to miss the opportunity to acknowledge the contribution of all the faculty members of the department for their kind assistance and cooperation during the development of our project.

Finally, we acknowledge our friends for their contribution in the completion of the project.

Date:

Sig. (s):

Name (s):Mohammad Areeb/Ayush Tripathi/ Aman Singh Bohra/Yash Vinayak Singh.

Roll No.:2200290100103/2200290100051/2200290100017/2200290100197

ABSTRACT

Crop diseases are a major problem for farmers all around the world as they result in significantly reduced yields and annual financial losses totaling billions of dollars. These pathogens often strike without warning, making early detection critical for agricultural stability and food security. To contribute to solving this global problem, we developed an automated crop disease detection system which can accurately identify pathologies by analyzing a simple photograph of a leaf. We have implemented a deep learning-based system to classify these diseases using advanced neural architectures. Specifically, we trained a Convolutional Neural Network (CNN) based model on the extensive PlantVillage Dataset, which contains a total of 61,486 high-resolution images across 38 distinct classes. Each class represents a specific type of disease, while also including a dedicated category for healthy leaves to ensure the model distinguishes between normal and infected plant tissue. We have carefully pre-processed and augmented the data before training to ensure the model generalizes well to new, unseen data. This is done to improve the accuracy of our model in unpredictable real-world scenarios where environmental conditions are rarely perfect. Steps like resizing images to a fixed resolution, random cropping, rotating, and varying brightness levels were implemented as part of these preprocessing and augmentation stages. This helps the system handle common variations in sunlight and camera angles encountered in the field. The model then visually analyzes the diseased portions on the leaf and displays the diagnostic results quickly on a streamlined user interface. The model has achieved a classification accuracy of 98.7% for images under controlled conditions during initial testing, though we acknowledge this may vary depending on external factors like background noise, lighting, or camera sensor quality.

Currently, the system is built using a standard CNN model as its foundation. However, the future plan is to build and train more complex and computationally efficient architectures using models like EfficientNet, ResNet, and Vision Transformers to achieve even higher precision and faster processing. We also plan to integrate the trained model into a dedicated mobile application to allow farmers to identify plant diseases directly in the field without needing specialized equipment or laboratory access. The proposed system is cost-effective, fast, easy-to-use, and highly scalable, offering a powerful tool that can potentially help in reducing massive crop losses and aiding farmers in securing their livelihoods.

TABLE OF CONTENTS

LIST OF FIGURES

LIST OF TABLES

LIST OF ABBREVIATIONS

Abbreviation	Full Form
AI	Artificial Intelligence
ML	Machine Learning
DL	Deep Learning
CNN	Convolutional Neural Network
RNN	Recurrent Neural Network
SVM	Support Vector Machine
KNN	K-Nearest Neighbors
FL	Federated Learning
GAN	Generative Adversarial Network
DC-GAN	Deep Convolutional Generative Adversarial Network
ViT	Vision Transformer
YOLO	You Only Look Once
RGB	Red, Green, Blue
GLCM	Gray-Level Co-occurrence Matrix

CHAPTER 1
INTRODUCTION

1.1 INTRODUCTION

Agriculture is very important for providing food to the increasing population around the world and is a key part of the economy in many countries. As the global population continues to grow rapidly, the demand for food is also increasing at a significant rate, making agriculture more important than ever before. This is especially important in developing nations like India, where a large part of the population depends on farming for their daily livelihood, either directly as farmers or indirectly through related sectors such as transportation, storage, and food processing. Agriculture is not just about food production; it is closely linked to food security, rural development, employment opportunities, and the overall economic stability of a country. When farming performs well, it ensures a steady supply of food, helps maintain stable prices in the market, and improves the standard of living for farmers and their families. However, when agriculture faces challenges, the impact can be severe, leading to food shortages, increased prices, unemployment, and economic instability both locally and globally.

Plant diseases are one of the major factors that negatively affect agricultural productivity. They significantly reduce both the quantity and quality of crops, which directly impacts farmers' income and food availability. Infected plants usually show visible symptoms such as leaf spots, yellowing or discoloration, wilting, stunted growth, and unusual patterns on leaves or fruits. In some cases, fruits may rot or leaves may dry out prematurely, further affecting plant health. These symptoms not only reduce the visual appeal of crops but also lower their nutritional value and overall quality. As a result, such crops become less desirable in the market, making it difficult for farmers to sell them at a good price. This leads to financial losses and sometimes even forces farmers to discard their produce completely, increasing wastage.

Apart from the reduction in yield and quality, plant diseases also increase the cost of farming. Farmers often need to spend extra money on pesticides, fertilizers, fungicides, and other chemical treatments to protect their crops and control the spread of diseases. However, a major challenge is that diseases are not always correctly identified. In many cases, farmers rely on their own experience or limited knowledge, which can lead to incorrect diagnosis. This results in the use of wrong chemicals or excessive application of pesticides. Such practices not only increase the overall cost of cultivation but also harm the environment by polluting soil and water resources. Overuse of chemicals can also affect human health and reduce soil fertility in the long term. According to research, plant diseases are responsible for causing approximately 20% to 40% of crop losses worldwide every year, which highlights the seriousness and widespread nature of this problem.

Another critical issue is the late detection of plant diseases. In many real-life situations, diseases are identified only after they have already spread significantly across the field. Since many plant diseases spread very quickly, even a small infection can turn into a major outbreak within a short period of time. Environmental factors such as humidity, temperature, and wind can further accelerate the spread of diseases. If the disease is not detected at an early stage, it becomes much more difficult to control, often leading to severe damage or even complete crop failure. This not only affects the income of farmers but also disrupts the supply of food in the market, which can result in shortages and higher prices for consumers.

Traditionally, plant disease detection has been done through manual inspection, where farmers or agricultural experts examine plants visually to identify signs of disease. While this method is simple and widely used, it has several limitations. It requires experience, knowledge, and expertise to correctly identify diseases, and even then, human errors can occur. Additionally, manual inspection is time-consuming and not practical for large-scale farms where thousands of plants need to be monitored regularly. The situation becomes even more challenging in rural or remote areas where access to agricultural experts is limited or unavailable. As a result, many farmers are unable to get timely and accurate advice, which further worsens the problem.

Because of all these challenges, there is a strong need for a better, faster, and more reliable solution for detecting crop diseases. This is where modern technologies like Artificial Intelligence (AI) and image-based systems play an important role. These technologies use advanced algorithms and machine learning models to analyze images of plant leaves and identify diseases automatically with high accuracy. Unlike traditional methods, AI-based systems do not rely on human judgment and can provide consistent and quick results. Farmers can use smartphones or cameras to capture images of crops and get instant feedback about the health of their plants.

By using such advanced systems, farmers can detect diseases at an early stage and take appropriate actions before the problem becomes severe. Early detection helps in reducing crop losses, minimizing the use of unnecessary chemicals, and improving overall productivity. Moreover, these technologies can be integrated into mobile applications, making them easily accessible even in remote areas. This not only saves time and effort but also empowers farmers with better decision-making tools. In the long run, the adoption of AI-based disease detection systems can promote sustainable farming practices, protect the environment, and ensure a stable and sufficient food supply for the growing population.

1.2 PROJECT DESCRIPTION

With the rapid development of modern technologies like Artificial Intelligence (AI) and Computer Vision (CV), the way problems are solved in agriculture has changed significantly. These technologies have introduced new and efficient methods for handling tasks that were traditionally done manually, such as monitoring crop health and detecting plant diseases. In earlier times, identifying plant diseases relied heavily on human observation and expertise, where farmers or agricultural specialists would examine plants visually and make decisions based on their experience. However, this process was often slow, inconsistent, and prone to errors. Now, with the help of AI-based systems, it has become possible to analyze images of plant leaves automatically and extract meaningful information with high accuracy. This not only reduces the dependency on manual labor but also improves the speed, consistency, and reliability of disease detection, making it more suitable for large-scale agricultural use.

Among the various techniques used in this field, deep learning has emerged as one of the most powerful and widely adopted approaches. In particular, Convolutional Neural Networks (CNNs) have gained immense popularity because they are specifically designed for handling image-related tasks. CNNs work by processing visual data in multiple layers, allowing them to learn patterns and features from images in a way that is somewhat similar to human vision, but far more consistent and precise. This makes them especially effective for applications like plant disease detection, where identifying small differences in texture, color, and patterns is very important. Their ability to learn complex visual representations has made them a preferred choice for many researchers and developers working in the field of agricultural technology.

One of the key advantages of Convolutional Neural Networks (CNNs) is their ability to automatically extract features from images without requiring manual intervention. In traditional machine learning approaches, features such as color, texture, edges, and shapes need to be manually defined by experts, which can be time-consuming and may not always capture all the important details. CNNs overcome this limitation by learning features directly from raw image data during the training process. They can identify complex and subtle patterns such as leaf spots, discoloration, texture irregularities, vein distortions, and disease-specific structures that may not be easily visible to the human eye. This ability to automatically learn and adapt makes CNNs more accurate, flexible, and reliable when it comes to detecting a wide variety of plant diseases under different conditions.

In this project, we have developed a deep learning-based system capable of identifying 38 different types of plant diseases, along with healthy leaves, by analyzing an input image of a plant leaf. The system is designed with simplicity and ease of use as a primary focus, ensuring that even non-technical users, such as farmers, can operate it without difficulty. A user can simply capture an image using a smartphone or upload an existing image, and the system processes the image and predicts the disease within a few seconds. This quick and user-friendly interaction makes the system highly practical and accessible for real-world agricultural applications, especially in rural areas where technical expertise may be limited.

For training the model, we utilized the PlantVillage dataset, which contains more than 61,000 images of plant leaves across different crops and disease categories. This dataset is widely used in the research community due to its size, quality, and diversity. It includes images of both healthy and diseased leaves, covering a wide range of plant species and disease types. The availability of such a large and well-structured dataset plays a crucial role in improving the model's performance, as it allows the system to learn a wide variety of patterns and generalize better to new, unseen data. A diverse dataset also helps in reducing overfitting, ensuring that the model performs well not just on training data but also in real-world scenarios.

To further improve the performance and reliability of the system, several image preprocessing and data augmentation techniques were applied. Preprocessing steps such as resizing ensure that all input images are consistent in size, while normalization helps stabilize the training process by scaling pixel values. Data augmentation techniques such as rotation, flipping, zooming, and brightness adjustments are used to artificially expand the dataset and introduce variations. These variations simulate real-world conditions such as changes in lighting, camera angles, background noise, and environmental factors. As a result, the model becomes more robust and capable of accurately detecting diseases even when images are captured in different field conditions.

Due to these enhancements, the system becomes faster, more accurate, and more reliable in its predictions. It can process images in real time and provide results almost instantly, which is highly beneficial in agricultural settings where quick decision-making is crucial. Early detection of plant diseases allows farmers to take timely action, such as applying the correct treatment or isolating affected plants, thereby preventing the spread of diseases and minimizing crop damage. This not only improves crop yield but also reduces financial losses and promotes better farm management practices.

The main goal of this project is to make advanced Artificial Intelligence (AI)-based tools easily accessible, practical, and beneficial for farmers, especially those operating in rural or resource-constrained environments. In many regions, farmers still rely heavily on traditional methods and manual inspection to detect crop diseases, which can often be inaccurate, time-consuming, and dependent on expert knowledge that may not always be available. By introducing an AI-driven solution that is simple to use and easily deployable, this project aims to bridge the gap between modern technological advancements and conventional agricultural practices.

The proposed system is designed to be integrated into mobile or web-based platforms, allowing farmers to access it conveniently using smartphones or basic digital devices. The user interface is kept intuitive and straightforward so that individuals without any technical background can operate the system effectively. Farmers can simply capture or upload an image of a crop leaf, and the system will analyze it and provide immediate feedback regarding the presence and type of disease. This ease of use ensures wider adoption and usability in real-world scenarios.

Furthermore, this approach promotes the democratization of AI technology by making it accessible to a broader population rather than limiting it to researchers or large-scale agricultural enterprises. By enabling early disease detection, the system helps farmers take timely corrective actions, reduce crop losses, minimize excessive pesticide use, and improve overall crop quality. Ultimately, this contributes to increased agricultural productivity, economic stability for farmers, and the development of more sustainable and environmentally friendly farming practices.

CHAPTER 2

LITERATURE REVIEW

2.1. Overview of Plant Disease Detection

Plant disease detection has become a major focus in modern agriculture because it directly impacts crop yield, quality, and the overall safety of the food supply. As the global population continues to grow, ensuring healthy and disease-free crops is more important than ever, since even small infections can lead to significant losses if not identified and treated on time. In the past, farmers mainly relied on manual inspection and expert knowledge to identify plant diseases, which was often time-consuming, inconsistent, and prone to human error. However, with the rapid advancement of technology, these traditional methods are gradually being replaced by smart and automated systems that can detect diseases more quickly, accurately, and on a much larger scale. These systems are especially useful for early detection, which is critical in preventing the spread of diseases and minimizing crop damage. Technologies such as image processing, machine learning, and deep learning have significantly improved the way plant health is monitored, as they enable computers to analyze images of plant leaves and identify patterns, symptoms, and irregularities that may not always be visible to the human eye. In particular, deep learning models can learn complex features from large datasets, making them highly effective in recognizing different types of plant diseases under varying conditions. As a result, automated disease detection systems are making agriculture more efficient, reliable, and scalable by reducing manual effort, minimizing errors, and providing fast and accurate results. Overall, these advancements are transforming traditional farming practices into smarter and more sustainable approaches, helping farmers improve productivity, reduce losses, and ensure a stable and safe food supply.

2.1.1 Need for Automated Disease Detection

In recent years, there has been a significant rise in research aimed at improving the detection of plant diseases through advanced computational technologies. This surge in interest is mainly due to the pressing need to minimize crop losses, boost agricultural productivity, and secure global food supplies for an ever-increasing population. Research shows that plant diseases cause a large portion of yearly crop loss around the world, making early and precise detection essential for sustainable farming practices.

Traditionally, identifying plant diseases has depended on manual checks conducted by agricultural specialists.

Although this method works in certain situations, it is slow, requires a lot of effort, and heavily relies on human knowledge. Additionally, such expertise is often lacking in rural and remote regions, where farmers encounter the most difficulties. Manual techniques are also susceptible to personal bias and variability, resulting in unreliable diagnoses. These shortcomings have led to a strong call for automated, dependable, and scalable alternatives.

2.1.2 AI-Based Techniques in Agriculture

To tackle these challenges, Artificial Intelligence (AI) has become a valuable tool in modern farming. Scientists are looking into various techniques like Machine Learning (ML), Deep Learning (DL), hyperspectral imaging, and federated learning (FL). These methods allow for automatic analysis of plant pictures, leading to quicker and more precise identification of diseases.

Alongside these methods, different types of learning approaches such as supervised, unsupervised, semi-supervised, and self-supervised learning are being studied.

Supervised learning works well when there are labeled datasets available, but semi-supervised and self-supervised learning help cut down on the need for large annotated data sets, which can be costly and take a lot of time to prepare. Some research focuses on achieving very high accuracy in controlled settings, while other efforts are aimed at creating affordable and reliable systems that can work effectively in real-world farming situations.

2.2. Traditional and Deep Learning Approaches

The methods used for detecting plant diseases have evolved significantly over time, shifting from traditional machine learning techniques to more advanced deep learning approaches. In earlier research, scientists relied heavily on manually designed features such as color, texture, and shape, along with classical algorithms like Support Vector Machines (SVM) and k-Nearest Neighbors (KNN). While these methods provided reasonable results, they had several limitations, including a lack of flexibility, dependence on expert knowledge for feature extraction, and difficulty in handling large and complex datasets. As computational power increased and large image datasets became more widely available, deep learning techniques began to gain popularity, with Convolutional Neural

Networks (CNNs) playing a central role. CNNs are capable of automatically learning features from raw images, which allows them to capture complex patterns and variations more effectively than traditional methods. This has led to significant improvements in both the accuracy and speed of plant disease detection systems. Furthermore, the introduction of transfer learning has made it possible to use pre-trained models, such as those trained on large-scale datasets, and adapt them to specific agricultural tasks with relatively less data and training time. This approach has been especially useful in scenarios where collecting large amounts of labeled agricultural data is challenging. Despite these advancements, there are still several challenges in real-world applications, such as variations in lighting, background noise, different environmental conditions, and the need for models to generalize well outside controlled datasets. Overall, this section discusses the progression from traditional machine learning to modern deep learning approaches, including CNN-based models and transfer learning, while also highlighting the practical challenges that still need to be addressed.

2.2.1 Traditional Machine Learning Methods

Traditional machine learning methods like Support Vector Machines (SVM) and K-Nearest Neighbors (KNN) were some of the first techniques used to detect plant diseases. These approaches depend on manually created features, where experts in the field design characteristics based on the color, texture, and shape of infected areas. Tools like Gray-Level Co-occurrence Matrix (GLCM) are often used to analyze texture.

While these methods can work well enough, their success is largely based on how well the features are engineered.

They usually struggle to perform well across different datasets and changing environmental factors. Plus, creating these features is a slow process that needs specialized knowledge, which makes these methods hard to scale up.

2.2.2 Convolutional Neural Networks (CNNs)

To get around the problems with older methods, scientists are now using deep learning, especially Convolutional Neural Networks (CNNs). These networks can automatically learn important features from images without needing people to pick them out by hand.

In CNNs, the early layers pick up basic things like edges and textures, while the deeper layers understand more complicated patterns linked to diseases.

This step-by-step learning helps CNNs do a great job at recognizing images, especially when there are many different plants and diseases involved. Because of this, CNNs are now the go-to method for detecting plant diseases.

2.2.3 Transfer Learning Techniques

A big step forward in this area is the use of transfer learning, where models like VGG16, InceptionV3, and ResNet50 that were already trained on big datasets like ImageNet are further adjusted for tasks like identifying plant diseases. These models can be used for specific farming purposes even when the data set is not very large.

In 2023, a study compared several deep learning models along with traditional methods such as SVM and KNN.

The results showed that the InceptionV3 model, which was trained using the Adam optimizer, reached an accuracy of about 98.01% on the PlantVillage dataset. This shows how effective deep learning and transfer learning are at recognizing complex visual patterns.

2.2.4 Limitations in Real-World Scenarios

Even though many plant disease detection models show high accuracy in controlled environments such as laboratory datasets, they often struggle when applied in real-world conditions. In actual farming scenarios, several unpredictable factors can affect performance, including variations in lighting (such as shadows or excessive sunlight), complex and cluttered backgrounds, partial occlusion of leaves by other objects, and differences in image quality due to camera type or user handling. Additionally, environmental conditions like dust, moisture, and uneven angles further complicate image capture. These factors can make it difficult for models to correctly identify disease patterns, leading to reduced accuracy and reliability. This gap between controlled performance and real-world effectiveness highlights the need for more robust, adaptable, and generalized models that can handle diverse conditions and still deliver consistent results in practical agricultural settings.

2.3. Privacy-Preserving and Distributed Learning

As agriculture increasingly adopts data-driven technologies, concerns about data privacy, security, and efficient data utilization are becoming more important. In real-world scenarios, agricultural data is often collected from multiple sources such as farms, drones, sensors, and IoT devices, making it difficult to gather and store all the data in a single centralized location. This distributed nature of data raises issues related to data ownership, privacy, and secure sharing, especially when sensitive or location-specific information is involved. To address these challenges, privacy-preserving techniques have gained significant attention, as they allow systems to learn from data without directly accessing or exposing it. One of the most promising approaches in this area is Federated Learning (FL), which enables multiple devices or organizations to collaboratively train machine learning models while keeping their data stored locally. Instead of sharing raw data, only model updates or parameters are exchanged, which helps maintain privacy and reduces the risk of data breaches. Federated Learning can also be effectively combined with advanced deep learning models to improve performance while still ensuring secure data handling. However, it also faces challenges such as communication overhead, differences in data distribution across devices, and system complexity. Despite these limitations, ongoing research is focused on improving its efficiency, scalability, and robustness, making it a powerful solution for real-world agricultural applications.

2.3.1 Need for Data Privacy

In real-world agricultural environments, data is continuously collected from a wide range of sources such as farms, soil sensors, weather stations, drones, and IoT-based monitoring devices. This data is extremely valuable because it contains important information about farm conditions, crop types, soil health, productivity levels, and environmental factors. However, much of this data is sensitive in nature, as it may reveal private details like the exact location of farms, farming practices, yield quantities, and business-related information. Sharing such data openly can raise serious privacy and security concerns for farmers and agricultural organizations. Additionally, in many rural areas, the lack of proper digital infrastructure and reliable connectivity makes it difficult to gather and store all this data in a centralized system. These limitations highlight the need for solutions that can utilize distributed data effectively while ensuring that sensitive information remains secure and private.

2.3.2 Federated Learning (FL)

Federated Learning (FL) has emerged as an effective solution to address the challenges of data privacy and decentralized data storage. In this approach, instead of sending raw data to a central server, the model is trained locally on individual devices such as smartphones, sensors, or farm-based systems. Each device processes its own data and updates the model based on local learning. Only the learned model parameters or updates are then shared with a central server, where they are combined to improve the overall model. This process allows multiple participants to collaborate in building a powerful machine learning model without ever sharing their actual data. As a result, data privacy is preserved, and the risk of data breaches is significantly reduced. Additionally, federated learning reduces the need for large-scale data transfer, making it more efficient and suitable for environments with limited bandwidth or connectivity.

2.3.3 FL with Advanced Models

In recent years, researchers have started integrating federated learning with advanced deep learning models such as Convolutional Neural Networks (CNNs) and Vision Transformers to further enhance performance. These hybrid approaches combine the strong feature extraction capabilities of CNNs with the global attention mechanisms of Vision Transformers, resulting in more accurate and robust models. When applied in a federated learning setup, these models can learn from diverse datasets distributed across different locations without compromising data privacy. Studies have shown that such combinations can achieve high accuracy levels comparable to centralized training methods, even when data is scattered and varies in quality. This makes them highly suitable for real-world agricultural applications, where data is often heterogeneous and distributed. However, challenges such as communication efficiency, model synchronization, and handling uneven data distribution still need to be addressed to fully realize their potential.

2.3.4 Challenges in Federated Learning

Even though it has benefits, federated learning brings some difficulties:

- High communication costs from constant data sharing
- Diverse data across different clients (non-IID distribution)
- Problems with keeping everything in sync because of different device abilities.

2.3.5 Optimization Techniques

To tackle these challenges, researchers are exploring advanced techniques such as model compression, adaptive aggregation, and asynchronous training to improve the efficiency and stability of distributed learning systems. Model compression focuses on reducing the size and complexity of machine learning models, making them easier to deploy on devices with limited computational power, such as smartphones or IoT devices used in agriculture. This not only speeds up processing but also reduces communication costs during training. Adaptive aggregation methods are designed to intelligently combine updates from different devices by considering factors like data quality, device reliability, and contribution significance, which helps in improving overall model performance. Additionally, asynchronous training allows devices to update the model at different times instead of waiting for all participants to synchronize, thereby reducing delays and making the system more flexible and scalable.

Together, these approaches help address key challenges in real-world implementations, such as limited resources, communication overhead, and inconsistent data distribution, ultimately making federated learning systems more practical and robust for agricultural applications.

2.4 Review Studies and Learning Paradigms

As research in plant disease detection continues to grow rapidly, many review studies have played an important role in summarizing existing approaches and identifying potential directions for future work. These studies provide a comprehensive overview of various machine learning and deep learning techniques used for disease detection, including methods such as Support Vector Machines (SVM), Convolutional Neural Networks (CNNs), and hybrid models. They highlight the strengths of these approaches, such as high accuracy and automation, while also discussing their limitations, including issues like overfitting, lack of generalization, and dependency on large labeled datasets. By comparing different models and datasets, review papers help researchers understand which techniques are most suitable for specific agricultural scenarios and where improvements are still needed.

In addition to traditional supervised learning approaches, newer learning paradigms such as semi-supervised and self-supervised learning are gaining increasing attention in recent years. These methods are particularly useful in situations where labeled data is limited or expensive to obtain, which is often the case in agriculture. Semi-supervised learning makes use of a small amount of labeled data along with a large amount of unlabeled data to improve model performance, while self-supervised learning allows models to learn useful representations from unlabeled data by creating their own learning tasks. These approaches reduce the dependency on manual annotation, which can be time-consuming and require expert knowledge. As a result, they make it easier to develop scalable and cost-effective disease detection systems.

Overall, recent review studies not only provide valuable insights into the progress made in plant disease detection but also highlight emerging trends and technologies that are shaping the future of this field. By focusing on modern learning methods and addressing current limitations, researchers are working towards building more robust, efficient, and adaptable systems. These advancements are expected to play a key role in the development of intelligent agricultural solutions that can operate effectively in real-world conditions, ultimately improving crop health monitoring and supporting sustainable farming practices.

2.4.1 Review of Existing Models

Several review papers have looked closely at how Artificial Intelligence and image processing methods have advanced in detecting plant diseases. These studies cover a variety of techniques, such as Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), Support Vector Machines (SVMs), and mixed models that use different methods together to boost results. CNNs are popular because they are good at finding spatial patterns in plant images, while RNNs are used when dealing with data that changes over time or shows a sequence of events. Hybrid models that mix traditional machine learning with deep learning try to use the best parts of both to get more accurate and reliable outcomes. These reviews not only show how performance has improved but also point out major challenges like overfitting, difficulty in applying results to new situations, and the need for large amounts of data.

2.4.2 Semi-Supervised and Self-Supervised Learning

Recent studies have increasingly focused on semi-supervised and self-supervised learning approaches to tackle the issue of limited labeled data in agricultural datasets. In real-world situations, gathering and labeling large amounts of plant disease images requires expert input, which makes the process both time-consuming and expensive. Semi-supervised learning uses a small amount of labeled data combined with a larger set of unlabeled data to enhance model performance. Meanwhile, self-supervised learning allows models to learn useful features by completing auxiliary tasks, like image reconstruction or contrastive learning, without the need for explicit labels. These methods greatly reduce the need for annotated data while still achieving strong accuracy. Because of this, they are seen as highly effective for agricultural applications, where data is often scarce and scalability is crucial.

2.5. Challenges in Real-World Deployment

Despite significant advancements in plant disease detection using Artificial Intelligence and deep learning techniques, deploying these models in real-world agricultural settings remains a complex and challenging task. While many approaches achieve high accuracy under laboratory conditions or on well-curated datasets, their performance often drops when applied in practical field environments. This is mainly because real-world conditions are far more unpredictable and diverse compared to controlled settings. Factors such as varying lighting conditions, shadows, weather changes, background clutter, and partial occlusion of leaves can significantly affect the quality of input images and, consequently, the model’s predictions. Additionally, data collected in real environments often comes from multiple sources, including smartphones, drones, and sensors, which introduces inconsistencies in image resolution, angles, and quality.

Another major challenge lies in computational constraints and scalability. Many deep learning models are resource-intensive and require high processing power, making them difficult to deploy on low-cost devices like smartphones or edge devices commonly used by farmers. Limited internet connectivity in rural areas further complicates cloud-based solutions, as real-time data transfer may not always be feasible. Moreover, models trained on specific datasets may struggle to generalize across different crops, regions, and environmental conditions, leading to reduced reliability in diverse agricultural scenarios. These challenges create a noticeable gap between research-level performance and real-world applicability.

Therefore, it becomes essential to focus not only on improving model accuracy but also on enhancing robustness, adaptability, and efficiency. Developing lightweight models, improving generalization through diverse datasets, and incorporating techniques that can handle environmental variability are key steps toward practical deployment. Addressing these issues will ensure that plant disease detection systems are not only accurate in theory but also reliable, scalable, and effective in real agricultural environments, ultimately benefiting farmers and supporting sustainable farming practices.

2.5.1 Generalization Issues

One of the most critical challenges in plant disease detection systems is their limited ability to generalize effectively in real-world conditions. While many deep learning models demonstrate high accuracy during training and testing phases, these evaluations are often conducted under controlled environments using curated datasets. Such datasets typically consist of images captured under ideal conditions, including uniform lighting, clean and uncluttered backgrounds, high-resolution imaging, and consistent camera positioning. As a result, the models tend to learn patterns that are highly specific to these controlled settings rather than robust, real-world features.

In practical agricultural environments, however, conditions are far more complex and unpredictable. Factors such as varying illumination (e.g., sunlight intensity, shadows, cloudy weather), occlusions from overlapping leaves or external objects, soil backgrounds, pest interference, and diverse camera angles significantly affect image quality. Additionally, images captured using mobile devices may introduce noise, blur, and compression artifacts. These variations create a domain gap between training data and real-world inputs, making it difficult for models to maintain consistent performance.

Consequently, models often suffer from overfitting, where they perform exceptionally well on training data but fail to adapt to unseen scenarios. This lack of generalization reduces their reliability and practical usability for farmers, especially in rural or resource-constrained settings. To address this issue, researchers are increasingly focusing on techniques such as data augmentation, transfer learning, domain randomization, and the inclusion of diverse field datasets during training. Improving generalization is essential for developing scalable, real-time disease detection systems that can operate effectively across different agricultural environments.

2.5.2 Cross-Dataset Limitations

Another significant challenge in plant disease detection is cross-dataset inconsistency, which refers to the inability of models trained on one dataset to perform effectively on other datasets. This issue becomes particularly evident when models trained on benchmark datasets, such as the widely used PlantVillage dataset, are tested on images collected from real-world farming environments or alternative data sources. Although these models may achieve very high accuracy on their original datasets, their performance often drops considerably when applied to new datasets.

The root cause of this limitation lies in the inherent differences between datasets. Variations in image acquisition conditions, including lighting, resolution, background complexity, and camera devices, can significantly alter the visual characteristics of the data. Moreover, datasets may differ in terms of crop varieties, disease types, geographical distribution, and growth stages of plants. These discrepancies result in different data distributions, making it challenging for models to extract features that are universally applicable.

As a result, models tend to learn dataset-specific features rather than generalized representations, limiting their transferability across domains. This poses a major obstacle for deploying plant disease detection systems at scale, where adaptability to diverse agricultural conditions is essential. To overcome cross-dataset limitations, several advanced approaches have been proposed, including domain adaptation, domain generalization, and transfer learning techniques. Additionally, creating large-scale, diverse, and representative datasets that capture real-world variability is crucial. Data augmentation strategies and synthetic data generation can also help bridge the gap between datasets. Addressing these challenges is key to building robust and flexible models capable of performing reliably across different environments and datasets.

2.6. Advanced Techniques in Plant Disease Detection

As artificial intelligence, deep learning, and imaging technologies continue to advance, the approach to plant disease detection is rapidly evolving beyond traditional machine learning methods. Earlier techniques relied on manually engineered features and simpler models, which, although useful, often struggled to capture the complex patterns present in real-world plant images. Today, researchers are exploring more advanced and specialized methods that leverage deep neural networks, high-resolution imaging, and large-scale datasets to improve detection accuracy and reliability. These modern approaches enable earlier identification of diseases, which is crucial for preventing their spread and minimizing crop damage. They also support faster decision-making by providing near real-time analysis, helping farmers take timely and informed actions in the field.

These newer techniques make use of improved data collection methods, such as drone imagery, satellite data, and sensor-based monitoring, along with sophisticated model architectures like Convolutional Neural Networks (CNNs), Vision Transformers, and hybrid systems. Enhanced data preprocessing, augmentation, and optimization strategies further contribute to better model performance by allowing systems to adapt to variations in lighting, background, and environmental conditions. As a result, these methods overcome many limitations of older systems, offering higher accuracy, better generalization, and greater efficiency in detecting a wide range of plant diseases.

However, despite these advantages, the adoption of such advanced techniques also introduces new challenges. These systems often require higher computational resources, making them more expensive to develop and deploy, especially in resource-limited agricultural settings. Their increased complexity can make them harder to understand, maintain, and integrate into existing farming practices. Additionally, scaling these solutions for widespread use across different regions and crop types remains a significant hurdle due to variations in data, infrastructure, and environmental conditions. Therefore, while modern approaches represent a major step forward in plant disease detection, there is still a need to balance performance with cost, simplicity, and scalability to ensure their practical usability in real-world agriculture.

2.6.1 Hyperspectral Imaging

Hyperspectral imaging is a sophisticated method that takes picture data across a broad range of wavelengths, including those not visible to the human eye. Unlike regular color imaging, which uses red, green, and blue light, hyperspectral imaging gives detailed information about the different wavelengths of light that a plant reflects. This helps detect small changes in a plant's health that aren't obvious to the naked eye, making it very useful for finding diseases early and analyzing plant conditions accurately. However, even though it has many benefits, hyperspectral imaging has some real-world challenges. It needs special and costly equipment, creates a lot of data that requires big storage space, and involves complicated steps to process and understand the data. These issues make it hard to use in real-time or in places where resources are limited, especially in farming areas that don't have much money or technology.

2.6.2 Crop-Specific and GAN-Based Models

To boost accuracy, some scientists have created models that are specific to certain crops and diseases. These models often use advanced methods like Deep Convolutional Generative Adversarial Networks (DC-GANs) along with Convolutional Neural Networks (CNNs) to improve how well they extract features and classify data. GANs can also help by creating fake images, which is useful when there's not enough real data available. Even though these methods work well in their own area, they aren't very flexible or adaptable, so they don't perform as well when applied to different crops or various farming situations.

2.6.3 Transformer and Hybrid Models

Recent progress in deep learning has brought about transformer-based models and mixed architectures that combine CNNs with networks like EfficientNet or Vision Transformers. These models can effectively capture both detailed local features, such as textures, spots, and lesions, and broader global information, like the overall structure and patterns of a leaf. By combining these elements, the models produce better feature representations and achieve higher accuracy in classification. Some research has even shown accuracies reaching up to 97.6%. However, these models often need a lot of computational power and strong hardware, which can make them hard to use in real-world situations.

2.6.4 Real-Time Detection Models

Real-time detection models, like YOLOv8, have become popular because they are fast and efficient. These models are built for identifying objects in images and can deliver quick results, making predictions almost instantly. This feature makes them very useful in actual field situations, where farmers need fast answers to make timely decisions. However, even though they are fast, these models might not always be perfectly accurate, and their effectiveness can be influenced by different conditions and image quality in real-life settings.

2.7. Proposed Model Advantages

The proposed plant disease detection model is designed not only to achieve high classification performance but also to address the practical limitations observed in existing approaches. While many current models focus primarily on improving accuracy, they often overlook important real-world challenges such as varying environmental conditions, limited computational resources, and ease of use for farmers. This model aims to strike a balance between accuracy, efficiency, and real-world applicability, making it more suitable for deployment in actual agricultural settings. By considering these factors, the system ensures that it is not only technically strong but also practically useful for end users.

One of the key strengths of the proposed model lies in its improved architecture, which is optimized to deliver reliable predictions while maintaining computational efficiency. The model is designed to be lightweight, allowing it to run on devices with limited processing power such as smartphones or edge devices commonly used in rural areas. This makes it highly accessible and practical for farmers who may not have access to high-end computing resources. At the same time, the model maintains high accuracy by effectively learning and identifying important features such as leaf patterns, discoloration, and disease-specific textures.

Another important aspect of the model is its simplicity and user-friendly design. The system is built in a way that requires minimal technical knowledge to operate, enabling farmers to easily use it in their daily activities. Users can simply capture or upload an image of a plant leaf, and the model quickly processes it to provide a prediction. This ease of use reduces dependency on experts and allows farmers to make faster decisions regarding crop health and treatment.

Additionally, the model is developed to perform well under a variety of real-world farming conditions. It is trained using diverse data and enhanced with preprocessing and augmentation techniques to handle variations in lighting, background, and image quality. This improves its robustness and ensures consistent performance even in challenging environments. Overall, the proposed system combines strong technical performance with practical usability, making it a reliable and effective solution for plant disease detection in real agricultural scenarios.

2.7.1 Efficiency and Accuracy

The proposed CNN-based model achieves a high classification accuracy of about 98.7%, showing its effectiveness in identifying various plant diseases. Unlike many complex models, this one has a balanced design that ensures accuracy without sacrificing computational efficiency. By carefully optimizing the network structure and training process, the model can extract important features without needing too many parameters or deep layers. This helps reduce training time and computational resources while still delivering strong performance, making it ideal for real-world use.

2.7.2 Practical Deployment

One of the main benefits of this model is its strong ability to perform well in real-life agricultural conditions, not just in controlled environments. Unlike many advanced systems that rely on expensive hardware or specialized imaging techniques, this model works using standard RGB images, which can be easily captured using regular smartphone cameras. This eliminates the need for costly equipment such as hyperspectral sensors or high-end imaging devices, making the system much more affordable and practical for everyday use. The data preprocessing techniques used in this model are simple yet effective, ensuring that the input images are standardized without requiring complex or time-consuming setup procedures. This simplicity allows the model to run smoothly even on devices with limited computational power. Additionally, the system is designed to be scalable, meaning it can be deployed across farms of different sizes, from small-scale farms to larger agricultural operations. Its low-cost implementation and minimal infrastructure requirements make it especially suitable for developing regions, where access to advanced technology may be limited.

2.7.3 Usability and Accessibility

The system is developed with a strong focus on usability and accessibility, ensuring that it can be easily adopted by farmers with little to no technical background. It can be integrated into a mobile application, allowing users to simply capture images of plant leaves using their smartphones and receive quick predictions about possible diseases. This real-time feedback enables farmers to take immediate action, such as applying the correct treatment or isolating affected plants, which helps prevent the spread of disease and reduces potential crop losses. The user interface is designed to be simple and intuitive, so even individuals who are not familiar with advanced technology can use the system without difficulty. This is particularly important in rural and remote areas, where access to expert agricultural advice and modern farming tools is often limited. By providing an easy-to-use and widely accessible solution, the system empowers farmers to monitor crop health independently, improve productivity, and make better-informed decisions for sustainable farming practices.

2.8. Conclusion of Literature Review

This method successfully combines accuracy, efficiency, cost-effectiveness, and ease of use, making it a well-balanced solution for modern agricultural challenges. In many existing plant disease detection systems, the primary focus is often placed on achieving high accuracy under controlled conditions. While accuracy is undoubtedly important, such an approach can overlook other critical factors required for real-world deployment, such as computational efficiency, affordability, and user accessibility. The proposed system addresses these gaps by ensuring that strong performance does not come at the expense of practicality. It is designed not only to deliver reliable predictions but also to function effectively in the diverse and often unpredictable conditions found in real farming environments.

One of the key strengths of this method lies in its efficiency. The model is optimized to process images quickly and provide near real-time results, which is essential in agriculture where timely decision-making can significantly impact crop health and yield. Early detection of plant diseases allows farmers to take immediate corrective actions, preventing the spread of infections and reducing potential losses. Unlike heavier and more complex models that require high-end hardware and long processing times, this system is lightweight and capable of running on commonly available devices such as smartphones or low-cost computing systems. This ensures that even farmers in resource-constrained settings can benefit from advanced technological solutions without needing expensive infrastructure.

Cost-effectiveness is another major advantage of this approach. Traditional disease detection methods may involve expert consultation, laboratory testing, or repeated use of chemical treatments, all of which can be expensive and time-consuming. In contrast, this system minimizes such dependencies by providing instant analysis based on simple image inputs. By reducing the need for excessive pesticide use through accurate diagnosis, it also helps lower input costs and promotes environmentally friendly farming practices. Furthermore, the use of widely available tools such as mobile cameras makes the system accessible to a larger population of farmers, including those in rural and remote areas.

Ease of use is equally important for the successful adoption of any technological solution in agriculture. The proposed system is designed with a user-friendly interface that requires minimal technical knowledge to operate. Farmers can simply capture or upload an image of a plant leaf, and the system automatically analyzes it to identify possible diseases. The simplicity of this process ensures that users do not need specialized training, making the technology more inclusive and practical. This is particularly beneficial in regions where access to agricultural experts is limited, allowing farmers to independently monitor and manage crop health.

Another significant aspect of this method is its adaptability to real-world conditions. The model has been trained and optimized using various preprocessing and data augmentation techniques, enabling it to handle differences in lighting, background noise, image quality, and environmental factors. This improves its robustness and ensures consistent performance even outside controlled datasets. As a result, the system is better equipped to deal with the challenges commonly faced in field conditions, such as varying weather, uneven lighting, and diverse crop environments.

Unlike many existing models that prioritize theoretical performance, this system emphasizes practical usability and scalability. Its balanced design makes it suitable for integration into mobile applications or web-based platforms, allowing for large-scale deployment across different regions and farming communities. By focusing on real-world applicability, the system bridges the gap between research and implementation, making it a viable solution for modern agriculture.

Overall, this method represents a comprehensive approach to plant disease detection by combining technical strength with practical considerations. It not only achieves high accuracy but also ensures efficiency, affordability, and ease of use, which are essential for real-world adoption. By addressing the limitations of existing systems and focusing on user needs, it stands out as a reliable and scalable solution. This makes it a strong candidate for widespread use in today's agriculture industry, where smart, efficient, and accessible technologies are needed to support sustainable farming and improve crop productivity.

CHAPTER 3

PROPOSED METHODOLOGY

Crop disease detection has always been carried out using conventional ways, but these ways have their limitations. In those approaches, crop disease detection is primarily based on farmers manually checking the crops or hiring professionals which tends to be tedious, inaccurate, time-consuming and costly when dealing with large farms. Additionally, farmers don't always have access to experts who can guide them throughout the crop cycle. These limitations were taken into consideration during this project. This project aims to detect plant diseases as effectively as possible in terms of being time-efficient, accurate and user-friendly to all.

The overall workflow of the system is simple. As depicted in Fig. 1, allow the user to upload an image of the leaf, and the system will identify the disease plaguing the plant. The user does not have to manually go through each disease one by one. Since the image-processing system is used in this project, farmers can use the application in-field using their phones.

This system is built upon a customized Convolutional Neural Network architecture that is used for image classification. Convolution Neural Networks (CNN) are used extensively in Computer Vision applications because they have the ability to learn features from images. The model that is used here has been trained on PlantVillage dataset which has 38 classes, consisting of diseases and healthy leaf classes. This provides a good mix of classes for our model to learn from.

One of the main benefits of using a CNN is that it doesn’t need people to pick out features by hand. In older machine learning methods, you have to manually define things like color, shape, and texture, which can be hard and can lead to mistakes. But with CNNs, the model figures out these features on its own during training. It can spot important visual clues like leaf spots, blight signs, mosaic patterns, mildew, and color changes. These clues are important for telling different plant diseases apart, and the model can tell them apart very accurately.

To make sure the system works well both in controlled settings and real-world situations, several image preparation and enhancement techniques are used.

Real-world images can change a lot because of things like light, camera quality, angle, and background. To handle these changes, images are first made the same size so the model can process them the same way every time. The pixel values are also adjusted to make training faster. Plus, techniques like rotating, flipping, changing brightness, and adjusting colors are used. These steps help the model handle different kinds of images in real life and make it more reliable.

After the image is prepared, it goes through the CNN model, which finds the important features and sorts the image into a disease category.

The system also gives a confidence score with the prediction. This score shows how sure the model is about its answer, which makes the results more trustworthy. A higher score means the model is more certain about the result, helping users feel more confident in the output.

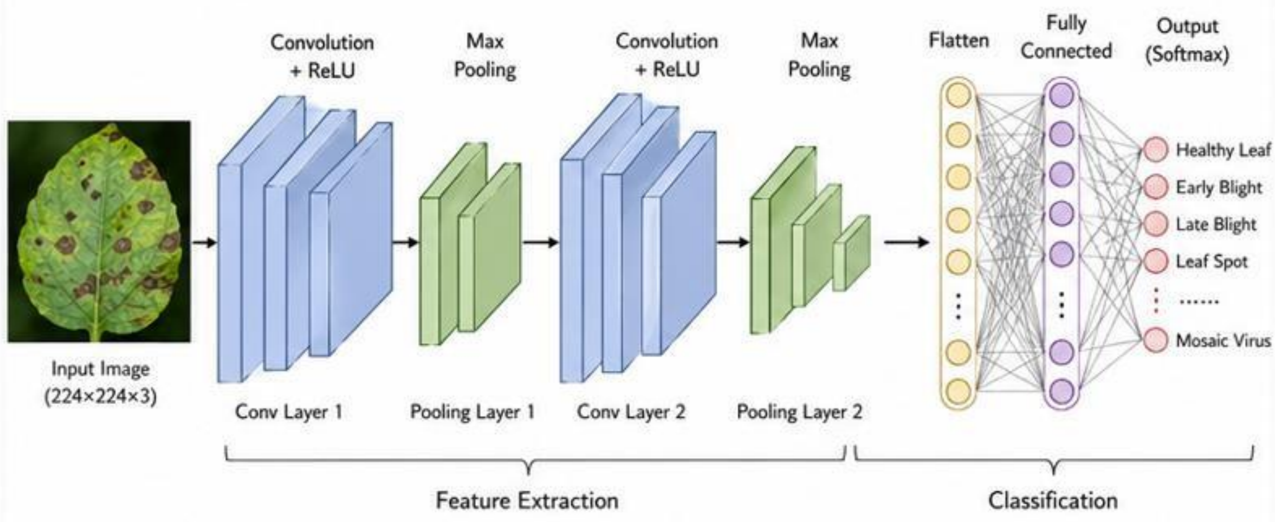

Fig. 1. General architecture of a Convolutional Neural Network (CNN).

The system is built to be easy to use through both web and mobile platforms. This makes it very practical because farmers can simply take a picture of a leaf with their smartphone and get instant results in just a few seconds. There's no need for special equipment or technical skills, which means the system is simple to use and can be accessed by many people.

This solution is especially helpful for farmers in faraway or under-resourced areas where it's hard to get help from agricultural experts, labs, or advanced tools.

By quickly and accurately identifying plant diseases, the system helps farmers act fast, like using the right treatment or taking steps to prevent the disease. This can greatly cut down on crop loss and boost farming output.

Besides detecting diseases, the system also has a visualization feature that makes it even more useful. It shows how sure it is about its predictions, keeps track of past disease findings, and helps farmers see how diseases change over time. This information can help farmers make smarter choices about how to take care of their crops. On a bigger scale, this data can also be used to monitor farming areas, track disease spread, and plan ways to stop diseases before they become a problem.

In total, this system offers a full and useful way to detect crop diseases by combining accuracy, speed, simplicity, and real-life usefulness.

It connects advanced technology with the everyday needs of farmers, making it a strong tool for modern agriculture.

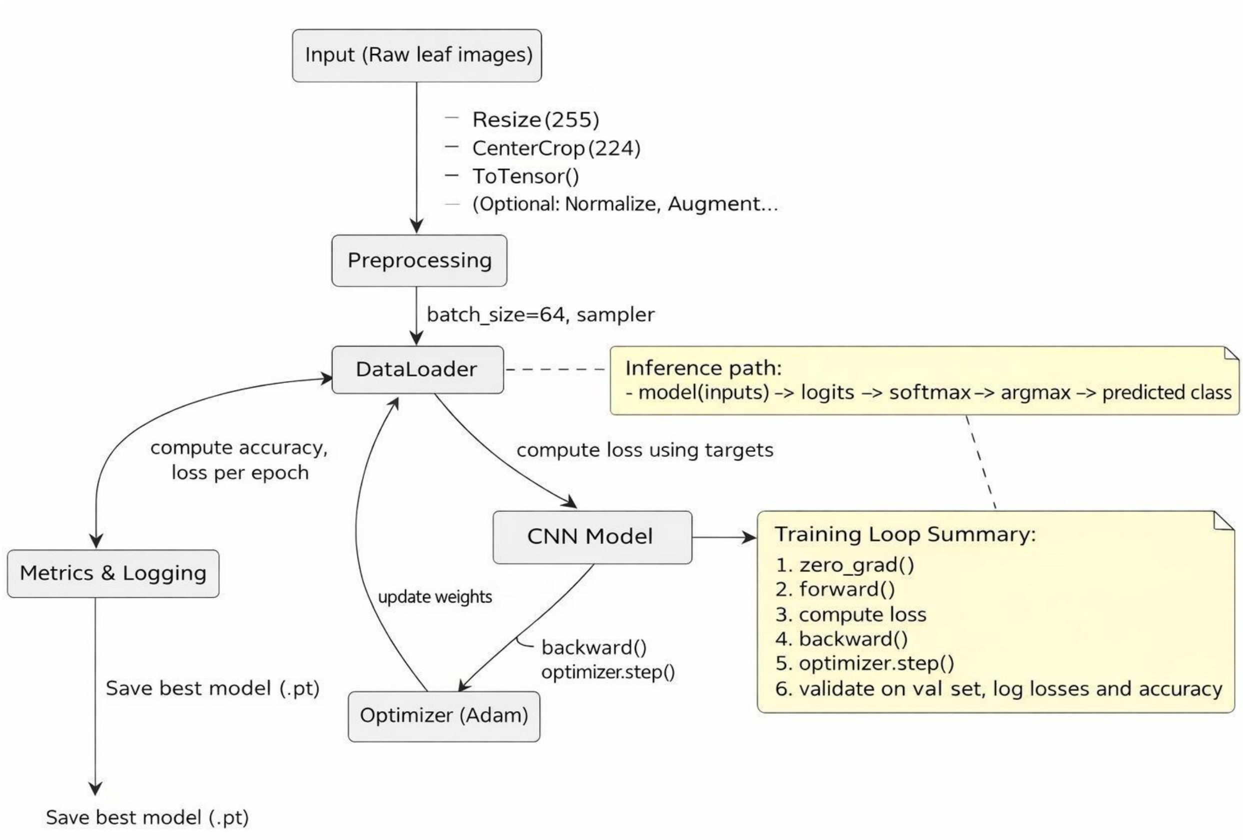

Proposed System

The proposed Crop Disease Detection System uses a modular deep learning approach that includes image preprocessing, CNN-based feature extraction, and real-time deployment. The system aims to provide an accurate, lightweight, and easy-to-use solution for real-world applications. Particular attention has been given to adapting the system for mobile use, allowing farmers to apply it directly in the field without needing costly hardware or advanced technical knowledge. By combining these elements, the system can effectively handle various crops and identify a broad range of diseases efficiently.

3.1. Requirement Analysis:

In traditional farming, people usually check plants by looking at them to see if they are sick. Farmers or experts look for signs like spots on leaves, changes in color, leaves drooping, or odd growth. This method has been used for a long time, but it has some problems. It takes a lot of time, especially on big farms where checking every plant is not easy. Also, it depends a lot on the experience of the person doing the checking. Even experienced people can make mistakes because many plant diseases look similar, which makes it hard to tell them apart.

Another issue is that there are not enough agricultural experts, especially in remote areas.

Farmers there might not have help from trained people, so they might use guesswork or advice from neighbors. This can lead to using wrong treatments, which can cost more and make things worse. Delaying the right diagnosis can let the disease spread and damage more crops.

To solve these issues, earlier methods used machine learning to automatically detect disease by looking at images of leaves.

These methods had some drawbacks. They mainly relied on features like color, texture, and shape, which were chosen by people. That meant the results depended a lot on how well those features were picked. If they weren't good, the system didn't work well.

Traditional machine learning models also performed well only in controlled settings.

They were trained on neat data, but they didn't handle real-world images well, like those with shadows, different lighting, or messy backgrounds. This made them less useful in real farming situations.

Hyperspectral imaging has been tested as a better solution.

It captures detailed light information beyond what humans can see, allowing early disease detection. However, it's expensive, needs special tools, and has complex data processing, making it hard for small farmers to use regularly.

To solve all these problems, this project proposes a lightweight system based on Convolutional Neural Networks (CNNs).

CNNs are powerful models that can automatically learn from images without needing manual feature selection. This makes the system more flexible, easier to use, and accurate.

The system is designed to work quickly and efficiently.

It's simple enough to run on smartphones, so farmers can take pictures of leaves in the field and get instant results. This real-time feature makes it very practical for everyday use.

Overall, this approach improves accuracy, speeds up the process, and makes disease detection more accessible and reliable.

Especially for farmers without access to experts or advanced tools, this system provides a helpful and easy-to-use solution.

3.2. System Design:

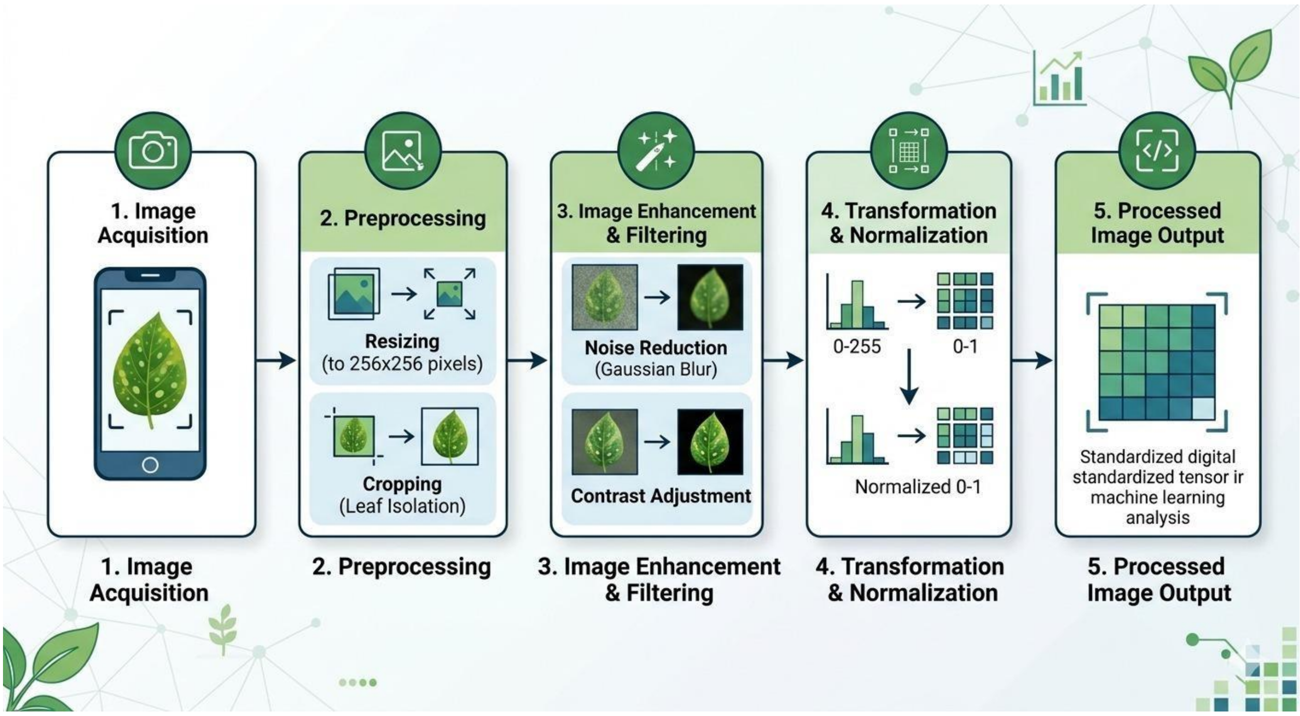

Image Processing

The system is built in a modular way and is split into three main parts, with each part handling a specific task in the full disease detection process. This modular approach makes the system easier to manage, understand, and expand for future upgrades. Each part works one after another to process the input image and produce reliable results.

3.2.1. Image Preprocessing Module:

The first step in the system is image preprocessing, which is very important for improving the model's performance. In real-world situations, images taken by farmers using mobile phones can be very different from each other. These differences might include changes in image size, lighting, camera quality, the angle of the shot, background noise, and shadows. If these variations are not handled well, they can hurt the model's accuracy.

To fix this, several preprocessing techniques are used before the images go into the CNN model. First, all images are resized to the same dimensions so the model gets a uniform input. Then, normalization is applied to adjust the pixel values to a standard range, which helps the model train faster and more steadily. Sometimes, noise reduction techniques are also used to remove unwanted disturbances from the image.

In addition, preprocessing may include simple cleaning steps like removing unnecessary parts of the image or focusing more on the leaf area.

These steps help highlight the important parts of the leaf while reducing unnecessary background details. Overall, this part of the process makes sure all input images are consistent, clean, and ready for the model to learn effectively.

3.2.2. Disease Classification Module (CNN):

This is the key part of the system where the real disease detection happens. A specially created Convolutional Neural Network (CNN) is used to automatically find features from the images and sort them into one of 38 disease types.

The CNN uses several layers, each with a different job:

- Convolution Operation:

In this step, the model uses many filters (or kernels) to look through the image and pick out important details like edges, textures, patterns, spots, and color changes.

These details help in recognizing the visual signs of different plant diseases.

- ReLU Activation:

After the convolution, the ReLU (Rectified Linear Unit) activation function is used.

It changes all negative numbers to zero and leaves the positive numbers as they are. This adds non-linearity to the model, allowing it to learn complicated patterns and connections in the data.

- Batch Normalization:

This layer makes the outputs from earlier layers more consistent, which helps in making the training process more stable and faster.

It also lowers the risk of overfitting and boosts the overall performance of the model.

- Max Pooling:

Max pooling is used to shrink the size of the feature maps while keeping the most important details. This helps in reducing the amount of computation needed and makes the model work better without losing crucial information.

- Fully Connected Layer:

Once the features are extracted, the data is flattened and sent through fully connected layers. These layers mix all the learned features together and help in making the final decision about the disease.

- SoftMax Layer:

The SoftMax function turns the output into probability values for each of the 38 classes. This makes it simpler to see which disease has the highest chance of being correct.

- Cross-Entropy Loss:

During training, this loss function measures how far the predicted output is from the correct label. The model uses this error to change its weights and get more accurate over time.

In total, this part allows the system to learn and recognize disease patterns on its own, without anyone having to manually pick out the features.

3.2.3. Deployment and User Interface:

After training, the model is turned into a real-time app that can be used on mobile phones and web browsers. The main goal of this part is to make the system simple and easy to use, so even people without technical skills, like farmers, can use it without trouble.

Users can just open the app, take a picture of a plant leaf with their phone camera, or upload a photo they already have.

The system quickly processes the image and gives the results in just a few seconds. This fast response makes the system very useful for use in the field.

The system gives the following outputs:

- The name of the predicted disease
- A confidence score that shows how certain the model is about its guess

In addition to these basic features, the system also has extra tools like saving past results and showing data through dashboards.

These dashboards let users watch how diseases happen over time, see patterns, and make smarter farming choices.

For example, a farmer can keep track of how often a certain disease appears and take steps to stop it from spreading.

On a bigger level, this data can help with monitoring crops, doing research, and planning better farming

strategies.
All together, the three parts of the system preprocessing, classification, and deployment create a full and efficient way to detect crop diseases.
The system is not only accurate but also easy to use, practical, and can grow with the needs of real-world farming situations.

3.3. Datasets:

For building and evaluating the proposed crop disease detection system, two types of datasets were considered: one for training and testing the model under controlled conditions, and another for future evaluation in real-world scenarios.

3.3.1. PlantVillage Dataset:

The main set of images used in this project is the PlantVillage dataset, which is commonly used in studies about detecting plant diseases. This dataset includes a total of 61,486 RGB images of plant leaves, divided into 38 different categories. These categories cover various types of plant diseases affecting different crops, as well as images of healthy leaves. The variety in this dataset allows the model to learn a wide range of disease patterns and enhances its overall ability to classify diseases accurately.
Most of the images in the PlantVillage dataset were taken under controlled conditions, meaning they have clear backgrounds, consistent lighting, and less noise.
This makes it easier for the model to identify important features while training. However, this also means the dataset does not fully reflect real-world situations, where images can be more complicated and varied.
To properly train and test the model, the dataset is split into three parts:
- 70% is used for training, where the model learns to recognize patterns from the data.
- 15% is used for validation, helping to adjust the model and track its performance during training.
- 15% is reserved for testing, where the model’s final performance is assessed using data it hasn't seen before.

This organized split helps prevent the model from becoming too focused on specific examples and ensures it can perform well on new, unseen data.

3.3.2. Field Dataset (Future Extension):

In addition to the PlantVillage dataset, another set of images from actual farms is being considered for future work. These images are taken directly from fields, making them more difficult and realistic compared to the PlantVillage dataset.

These images often have different problems such as:

- Poor lighting, like bright sunlight or dark shadows
- Background distractions like soil, other plants, or random objects
- Blurry or noisy images because of low-quality cameras or movement
- Leaves shown from different angles or positions

Because of these issues, using this field dataset helps test how well the model works in real situations. It also checks how well the model can adapt to new and different data.

Even though this dataset isn't fully used now, it's planned to be added in future system updates. Including real-world images makes the model stronger, more dependable, and better for use in real farms.

Overall, combining a structured dataset like PlantVillage with a more complex field dataset creates a solid base for building a useful and scalable system for detecting crop diseases.

3.4. Feature Extraction:

One of the biggest benefits of using Convolutional Neural Networks (CNNs) is their ability to automatically learn features from images at different levels. Unlike traditional machine learning methods, where you have to manually define features like color, texture, or shape, CNNs can learn these on their own during training. This makes them more powerful, flexible, and accurate for tasks like detecting plant diseases.

The way CNNs extract features is step by step.

The model starts by learning simple features in the earlier layers and then moves on to more complex and meaningful patterns in the deeper layers. This step-by-step learning helps the model understand the image in more detail and in a structured way.

Low-level features are the simplest parts:

In the first layers of the network, the CNN looks for basic elements like edges, lines, and color changes. These are the building blocks of any image. For instance, the model might find the edges of a leaf, changes in color, or simple patterns. Even though these are simple, they are important because they form the base for more complex features later.

Mid-level features come next:

As the data moves deeper into the network, the model starts to combine the simple features it found earlier. At this stage, it can identify things like spots, lesions, patches, and textures on a leaf. These features are directly linked to visible signs of plant disease and help in telling the difference between healthy and sick leaves.

High-level features are the most complex:

In the deeper layers of the network, the CNN learns abstract patterns that represent specific disease traits. These could include the overall shape of an infection, unique patterns of a disease, or combinations of several symptoms.

By learning features at each of these levels, the CNN gets a full picture of the image.

This lets the system accurately identify various plant diseases, like fungal infections, color changes, pest damage, and viral issues. The ability to automatically find and combine these features is why CNNs are so good at detecting crop diseases.

Overall, this step-by-step feature extraction helps the model move from simple visual details to complex disease patterns.

This makes the system accurate and reliable for identifying plant diseases in real-world situations.

3.5. Model Implementation:

In this project, the model is carefully designed to achieve a well-balanced combination of high accuracy, computational efficiency, and practical usability in real-world agricultural environments. Instead of focusing only on maximizing performance metrics, the approach emphasizes creating a system that can be realistically deployed and used by farmers and agricultural practitioners. The overall plan involves three key components: designing a custom Convolutional Neural Network (CNN) architecture specifically tailored for crop disease detection, comparing its performance with well-established transfer learning models, and selecting an effective training configuration that ensures optimal learning and stability. By integrating these elements, the system aims to provide a reliable and scalable solution that can operate efficiently even in resource-constrained environments such as rural areas where access to advanced computing infrastructure is limited.The core component of the system is a custom-designed Convolutional Neural Network (CNN) that is specifically built for identifying plant diseases from leaf images. The model is structured in a sequential and hierarchical manner, allowing it to learn visual features step by step, starting from basic patterns and gradually moving toward more complex disease-specific characteristics. The first major building blocks of the network consist of repeated groups of convolutional layers followed by Rectified Linear Unit (ReLU) activation functions and batch normalization layers. The convolutional layers play a crucial role in extracting important visual features from the input images, such as edges, textures, color gradients, and patterns that indicate the presence of disease. The ReLU activation function introduces non-linearity into the model, enabling it to capture complex relationships between these features and improve its learning capacity. Batch normalization further enhances the training process by stabilizing the distribution of inputs across layers, accelerating convergence, and reducing the risk of overfitting.

Following each group of convolutional operations, max pooling layers are applied to reduce the spatial dimensions of the feature maps. This process helps in decreasing computational complexity while retaining the most significant features extracted by the convolutional layers. By focusing only on the most important information, the model becomes more efficient and less sensitive to minor variations in the input images. In addition to this, dropout layers are incorporated into the architecture as a regularization technique to prevent overfitting. During training, dropout randomly disables a fraction of neurons, forcing the network to learn more generalized and distributed representations instead of relying on specific features. This improves the model's ability to perform well on unseen data and enhances its robustness in real-world scenarios where input conditions may vary significantly.

After the feature extraction phase, the model transitions into the classification stage. The learned feature maps are first passed through a flatten layer, which converts the multi-dimensional data into a one-dimensional vector. This flattened representation is then fed into a fully connected dense layer, where all the extracted features are combined to form a comprehensive understanding of the input image. The dense layer acts as a decision-making component, analyzing the relationships between features to determine the most likely disease category. Finally, the output layer uses the SoftMax activation function to generate probability scores for each of the 38 possible classes, including both diseased and healthy categories. The class with the highest probability is selected as the final prediction. This structured and efficient design ensures that the custom CNN model is both powerful and practical, making it highly suitable for real-world agricultural applications.

To evaluate the effectiveness of the proposed CNN model, a comparative analysis is conducted using several well-known transfer learning models that are widely used in image classification tasks. These models include MobileNetV2, ResNet50, and EfficientNet-B0, each of which has its own strengths and characteristics. MobileNetV2 is designed to be lightweight and computationally efficient, making it ideal for deployment on mobile devices and embedded systems. It uses depthwise separable convolutions to reduce the number of parameters and computational cost, enabling fast real-time performance. However, due to its relatively simpler structure, it may sometimes achieve slightly lower accuracy compared to deeper networks. ResNet50, on the other hand, is a deep neural network that utilizes residual connections to address the vanishing gradient problem, allowing it to learn very complex features effectively. While it offers high accuracy and strong feature extraction capabilities, it requires significantly more computational resources, which can limit its usability in low-resource environments. EfficientNet-B0 represents a more balanced approach by scaling network depth, width, and resolution in a systematic way, achieving a good trade-off between accuracy and efficiency. By comparing the custom CNN model with these advanced architectures, the study demonstrates that the proposed approach successfully achieves competitive accuracy while maintaining lower computational complexity, making it more practical for real-world deployment.

The training configuration of the model is another critical factor that contributes to its performance and stability. Careful selection of training parameters ensures that the model learns effectively without overfitting or underfitting. The Adam optimizer is used in this study due to its ability to combine the advantages of momentum-based optimization and adaptive learning rates. This allows the model to converge faster and handle complex optimization landscapes more efficiently. A learning rate of 0.001 is chosen to strike a balance between learning speed and stability; a higher learning rate could cause the model to overshoot optimal solutions, while a lower rate might slow down the training process. The batch size is set to 32, meaning that the model processes 32 images at a time during each training step. This batch size provides a good compromise between computational efficiency and stable gradient updates. The model is trained for 50 epochs, allowing it to pass through the entire dataset multiple times and learn the underlying patterns thoroughly without excessive memorization.

To measure the model's performance during training, the cross-entropy loss function is used. This loss function is particularly suitable for multi-class classification problems, as it quantifies the difference between the predicted probability distribution and the actual class labels. By minimizing this loss, the model continuously improves its predictions and becomes more accurate over time. The combination of an effective optimizer, well-chosen hyperparameters, and an appropriate loss function ensures that the training process is stable, efficient, and capable of producing a high-performing model.

Overall, the integration of a carefully designed CNN architecture, a thorough comparison with state-of-the-art transfer learning models, and a well-optimized training configuration results in a system that is both accurate and efficient. The model not only achieves strong performance in terms of disease classification but also remains practical for real-world use due to its lightweight design and scalability. This balanced approach ensures that the system can be effectively deployed in agricultural settings, providing farmers with a reliable tool for early disease detection and improved crop management.

3.6. Working (Pseudocode Explanation):

The proposed plant disease detection system is designed to operate in a clear, structured, and step-by-step manner, ensuring that each stage contributes to accurate, fast, and reliable predictions. The overall workflow is carefully organized so that even complex deep learning operations are executed seamlessly in the background while presenting a simple and user-friendly experience to the end user. From the moment an image is provided to the system until the final prediction is displayed, each component plays a critical role in maintaining efficiency, robustness, and usability. This structured pipeline not only improves the accuracy of the model but also ensures that the system can be easily deployed in real-world agricultural environments where quick and dependable results are essential.
The process begins with the model setup phase, where the system initializes by loading the pre-trained Convolutional Neural Network (CNN) model along with a predefined list of disease labels. These labels correspond to 38 different categories of plant leaves, including both healthy leaves and various disease types. This step is crucial because it prepares the system for inference by ensuring that all necessary components, including trained weights and classification mappings, are readily available. By preloading the model, the system reduces processing delays and ensures that predictions can be generated quickly once an image is provided. This initialization phase essentially sets the foundation for the entire detection process, making sure that the system is ready to handle new inputs efficiently and accurately.Once the system is initialized, the next step involves image acquisition, where the user provides an input image. The system is designed to be flexible and user-friendly, allowing users to either capture a new photo using a smartphone camera or upload an existing image from their device. This flexibility makes the system highly practical for real-world use, as farmers can use it directly in the field or analyze previously captured images at a later time. The ability to work with different input sources ensures that the system can adapt to various usage scenarios, making it accessible to a wide range of users with different levels of technical expertise.After the image is uploaded, it undergoes an important preprocessing stage to prepare it for analysis by the CNN model. During this phase, the image is resized to a fixed dimension of 256×256 pixels to match the input requirements of the model. Standardizing the image size ensures consistency across all inputs, which is essential for maintaining

stable model performance. In addition to resizing, the pixel values of the image are normalized, typically by scaling them to a standard range. This normalization process helps the model interpret the image data more effectively and improves convergence during prediction. In some cases, additional preprocessing steps such as noise reduction, contrast enhancement, or minor corrections may be applied to improve image quality. These enhancements help remove irrelevant variations and ensure that the model focuses only on meaningful visual information, thereby improving prediction accuracy.Following preprocessing, the image is passed into the CNN model for feature extraction. This is one of the most critical stages of the pipeline, where the model analyzes the image to identify important patterns and characteristics. The convolutional layers scan the image using filters to detect low-level features such as edges, lines, and color gradients. As the data moves deeper into the network, higher-level features such as textures, shapes, and disease-specific patterns are extracted. Activation functions and normalization layers further refine these features, ensuring that the most relevant information is highlighted. This hierarchical feature extraction process allows the model to build a detailed understanding of the image, enabling it to distinguish between different disease types with high accuracy.

Once the features are extracted, the system moves to the feature aggregation and classification stage. In this phase, the extracted features are combined and transformed into a format suitable for decision-making. The feature maps are flattened into a one-dimensional vector, which is then passed through fully connected (dense) layers. These layers analyze the relationships between different features and integrate them to form a comprehensive representation of the image. This step is essential because it translates the visual patterns identified earlier into meaningful information that can be used for classification. The dense layers act as the decision-making component of the model, preparing the data for the final prediction.The next stage involves generating probability scores using the SoftMax output layer. This layer converts the model’s outputs into a probability distribution across all 38 possible classes. Each class is assigned a probability value that indicates how likely it is that the input image belongs to that particular category. The SoftMax function ensures that all probability values sum up to one, making it easier to interpret the results. This probabilistic approach provides a clear understanding of the model’s confidence in each possible prediction, rather than giving a single rigid output.After calculating the probability scores, the system selects the class with the highest probability as the final prediction. This step ensures that the most likely disease (or healthy category) is chosen based on the model’s analysis. By relying on probability values, the system minimizes uncertainty and provides the most accurate classification possible. This decision-making process is both efficient and reliable, allowing the system to deliver consistent results across different inputs.Finally, the system presents the prediction results to the user in a clear and understandable format. The output includes the name of the detected disease along with a confidence score, which indicates how certain the model is about its prediction. This information is displayed in a user-friendly manner, enabling farmers or users to quickly interpret the results and take appropriate action. The inclusion of a confidence score is particularly useful, as it helps users assess the reliability of the prediction and decide whether further verification or immediate action is needed. By providing fast, accurate, and easy-to-understand results, the system empowers users to make informed decisions, ultimately supporting better crop management and reducing the impact of plant diseases.Overall, this step-by-step workflow ensures that the proposed system is not only technically sound but also highly practical for real-world agricultural applications. Each stage, from model initialization to result display, is optimized to maintain accuracy, efficiency, and usability. This structured approach makes the system a reliable tool for early disease detection and contributes to the advancement of smart and sustainable farming practices.The proposed system for detecting plant diseases works in a clear, step-by-step way to look at an image and give a good prediction. Each part of the process helps the system work well and give dependable results quickly.

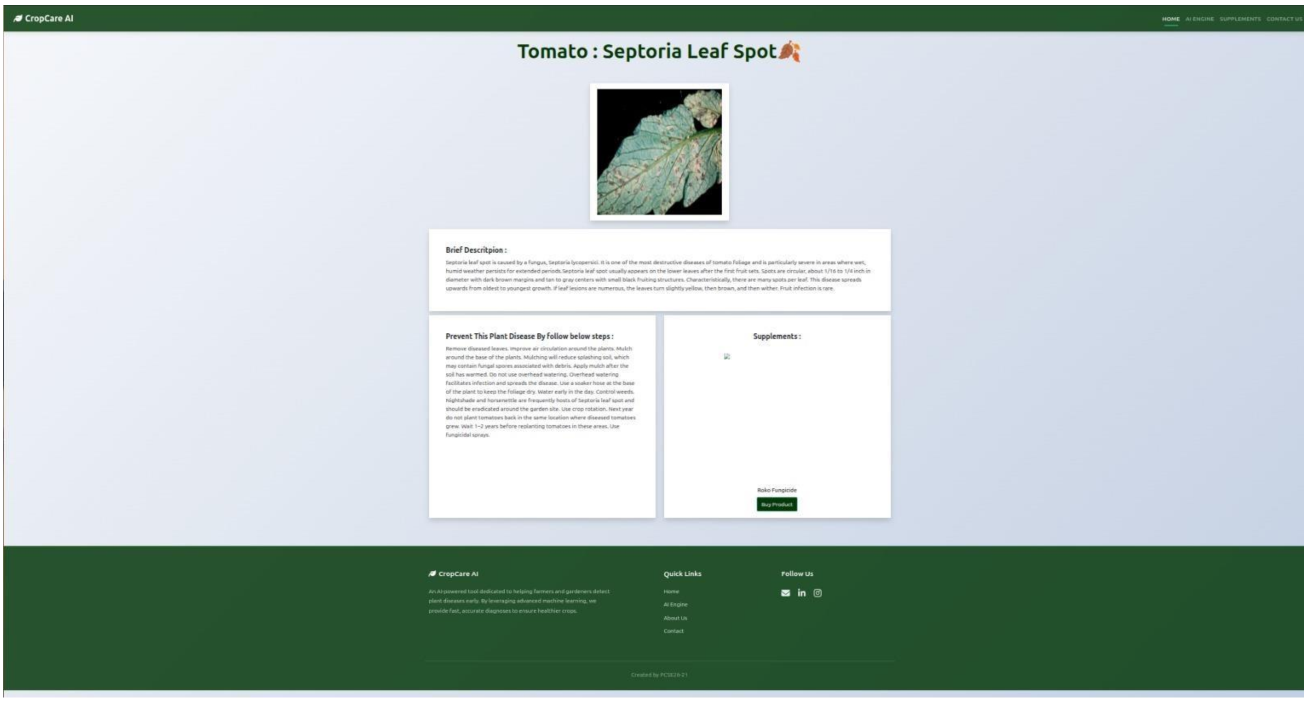

Project Image

CHAPTER 4

RESULTS AND DISCUSSION

This section presents a comprehensive and in-depth evaluation of the experimental results obtained from the proposed Convolutional Neural Network (CNN)-based crop disease detection system. The primary aim of this analysis is not only to measure the performance of the model in terms of accuracy but also to understand its learning behavior, stability during training, generalization capability on unseen data, and its practicality for real-world agricultural deployment. The evaluation focuses on three major aspects: the training performance of the model, its validation performance on unseen data (as illustrated in Fig.), and a comparative analysis with other widely used deep learning architectures (as shown in Table 1). The overall findings clearly demonstrate that the proposed model achieves a high level of accuracy while maintaining stability, efficiency, and robustness, making it highly suitable for real-life farming scenarios.

During the training phase, the model's learning progress was carefully monitored by tracking both training accuracy and validation accuracy over multiple epochs. These metrics provide valuable insights into how well the model is learning from the dataset and whether it can generalize its knowledge beyond the training data. It was observed that both training and validation accuracy increased gradually and consistently throughout the training process. There were no sudden spikes, drops, or irregular patterns in the accuracy curves, which indicates that the learning process was smooth and stable. This behavior is particularly important because unstable training can often lead to poor model performance or convergence issues. The steady improvement suggests that the model was effectively capturing meaningful patterns from the data rather than reacting to noise or random variations.

The accuracy curves for both training and validation datasets followed a similar upward trajectory and eventually reached a saturation point where further improvements became minimal. This phenomenon, known as convergence, indicates that the model has successfully learned the most important features required for classification and has reached its optimal performance level. A well-converged model is desirable because it ensures that computational resources are used efficiently and prevents unnecessary overtraining. Additionally, the smooth convergence of the model reflects the effectiveness of the chosen hyperparameters, such as learning rate, optimizer (e.g., Adam), batch size, and number of epochs. Proper tuning of these parameters plays a critical role in achieving optimal model performance. In terms of numerical results, the model achieved an exceptionally high training accuracy of 99.4% and a validation accuracy of 98.6%. These values indicate that the model performs extremely well on both seen and unseen data. The small difference of less than 1% between training and validation accuracy is a strong indicator of good generalization capability. This means that the model is not simply memorizing the training dataset but is instead learning generalized patterns that can be applied to new images. Generalization is one of the most important qualities of a machine learning model, especially in agricultural applications where real-world data can vary significantly in terms of lighting, background, orientation, and image quality.

Another critical component of the evaluation is the analysis of the loss function, which measures the error between predicted outputs and actual labels. Throughout the training process, both training loss and validation loss were observed to decrease steadily with each epoch. This continuous reduction in loss values indicates that the model is improving its predictions and becoming more accurate over time. A decreasing loss curve is a clear sign that the optimization process is working effectively. Furthermore, the difference between training loss and validation loss remained consistently small, which suggests that the model is not overfitting.

Overfitting is a common challenge in deep learning, where a model performs extremely well on training data but fails to generalize to new, unseen data. In such cases, the model essentially memorizes the training examples rather than learning meaningful patterns. However, in this study, the minimal gap between training and validation metrics confirms that overfitting has been successfully avoided. This is likely due to the use of effective regularization techniques, data augmentation strategies, and a well-designed model architecture. As a result, the model is able to maintain high performance across different datasets, making it more reliable for real-world applications.

In addition to performance evaluation, a comparative analysis was conducted between the proposed CNN model and other popular deep learning architectures such as MobileNetV2, ResNet50, and EfficientNet-B0. These models are known for their strong feature extraction capabilities and high accuracy in image classification tasks. However, they often come with increased computational complexity and require more processing power and memory. The proposed custom CNN model, on the other hand, achieves competitive accuracy while maintaining a lightweight structure. This makes it more suitable for deployment on devices with limited resources, such as smartphones or edge computing devices commonly used in agricultural settings. The balance between accuracy and efficiency is a key advantage of the proposed system, as it ensures both high performance and practical usability.

The experimental results also highlight the model's ability to effectively learn and identify important visual features associated with plant diseases. Using the PlantVillage dataset, which contains a diverse set of images across 38 classes, the model was able to detect various disease characteristics such as leaf spots, discoloration, texture irregularities, and complex disease patterns. The hierarchical feature extraction capability of CNNs allows the model to first capture low-level features like edges and color gradients, followed by higher-level features such as shapes and disease-specific patterns. This multi-level feature learning significantly enhances the model's classification accuracy and reliability.

Another important observation is the consistency of the model's performance across different disease categories. The model does not show bias toward specific classes and performs well across a wide range of plant diseases and healthy samples. This indicates that the model has been trained effectively on a balanced dataset and is capable of handling multi-class classification problems efficiently. Such consistency is crucial for real-world applications, where the model must be able to identify different types of diseases accurately.

From a practical perspective, the model's ability to deliver fast and accurate predictions makes it highly suitable for real-time applications. Farmers can use the system to quickly diagnose plant diseases by simply capturing images of leaves, enabling them to take immediate corrective actions. Early detection plays a vital role in preventing the spread of diseases, reducing crop losses, and improving overall productivity. The system's efficiency and responsiveness make it a valuable tool for modern precision agriculture.

In conclusion, the experimental results strongly validate the effectiveness of the proposed CNN-based crop disease detection system. The model demonstrates excellent accuracy, stable training behavior, strong generalization capability, and minimal overfitting. Its ability to perform well under diverse conditions, combined with its lightweight and efficient design, makes it a practical and scalable solution for real-world agricultural use. By bridging the gap between high-performance research models and real-world applicability, this system has the potential to significantly improve disease detection processes, support farmers in making informed decisions, and contribute to the advancement of smart and sustainable agriculture.

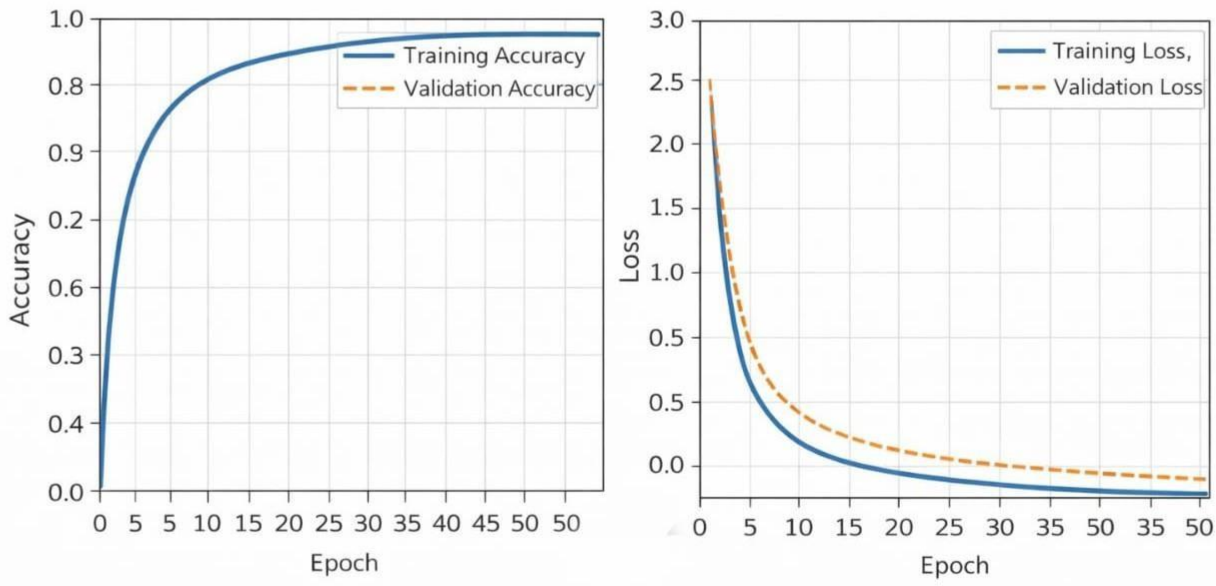

Training and Validation Accuracy of the proposed CNN model over 50 epochs.

Training and Validation Loss curves demonstrating stable convergence of the CNN

Training and Validation Accuracy & Loss Graph

Architecture	**Test Accuracy**	**Remarks**
Proposed CNN	98.7%	Achieved the highest accuracy among all evaluated models
MobileNetV2	96.3%	Optimized for low-resource and mobile-based applications
ResNet50	97.5%	Effective feature learning with minor overfitting tendencies
EfficientNet-B0	97.9%	Good trade-off between accuracy and computational cost

Table 1. Performance Evaluation of Different Deep Learning Architecture

CHAPTER 5

CONCLUSION AND FUTURE SCOPE

In this study, we conducted an extensive investigation into how modern deep learning techniques can be effectively utilized to detect crop diseases in a manner that is fast, accurate, reliable, and suitable for real-world agricultural environments. Agriculture plays a foundational role in sustaining human life by ensuring food security and supporting the livelihoods of millions of people, especially in developing countries where a large portion of the population depends on farming. One of the most critical challenges faced by farmers is the timely detection of plant diseases, as delayed identification can lead to severe crop damage, reduced yield, and significant financial losses. Recognizing this problem, the primary objective of this work was to design and develop an intelligent system that can automatically identify plant diseases at an early stage using image-based analysis. To achieve this, we proposed a deep learning-based solution built upon a custom-designed Convolutional Neural Network (CNN) architecture, specifically tailored for crop disease detection. The system is capable of analyzing images of plant leaves and accurately classifying them into different disease categories or identifying them as healthy, thereby enabling early intervention and better crop management.

The proposed model was trained using the widely recognized PlantVillage dataset, which contains a large and diverse collection of labeled images covering multiple crops and a wide range of disease types. This dataset serves as a strong foundation for training deep learning models due to its size, quality, and variety. It includes both healthy and diseased leaf images, allowing the model to learn distinguishing features effectively. The training process involved multiple stages, including data preprocessing, normalization, and augmentation, to ensure that the model could handle variations in image conditions. These steps helped improve the robustness of the model and reduced the chances of overfitting. To ensure a comprehensive evaluation, the performance of the proposed CNN model was compared with several well-established deep learning architectures, including MobileNetV2, ResNet50, and EfficientNet-B0. This comparative analysis was essential to assess not only the classification accuracy but also factors such as computational efficiency, scalability, and suitability for deployment in real-world agricultural environments. By evaluating multiple models under the same conditions, we were able to clearly understand the strengths and limitations of each approach.

The experimental results demonstrated that the proposed CNN model achieved the highest test accuracy of 98.7%, outperforming the other models considered in this study. This high level of accuracy indicates that the model is highly effective in identifying plant diseases based on leaf images. However, accuracy alone is not sufficient for practical applications; therefore, additional performance aspects such as generalization and consistency were also analyzed. The model exhibited strong generalization capability, meaning it performed well not only on the training dataset but also on unseen data. This is a crucial requirement for real-world applications, where input images can vary significantly in terms of lighting conditions, background complexity, orientation, and image quality. The ability of the model to maintain high performance under such variations highlights its robustness and reliability. Furthermore, the model showed consistent performance across multiple disease categories, ensuring that it does not favor specific classes while neglecting others. This balanced performance is particularly important in agricultural scenarios, where a wide variety of diseases may need to be detected accurately.

In comparison to other deep learning models, MobileNetV2 demonstrated slightly lower accuracy but proved to be highly efficient in terms of computational requirements. Its lightweight architecture makes it an excellent choice for mobile and edge-based applications, where processing power and memory are limited. On the other hand, ResNet50 and EfficientNet-B0 exhibited strong feature extraction capabilities and achieved competitive accuracy levels; however, they showed signs of overfitting when tested on real-world field images. This suggests that while these models perform well in controlled environments, their ability to generalize to real-world conditions may be limited. The proposed custom CNN model, in contrast, successfully balances accuracy and efficiency, making it more suitable for practical deployment. Its relatively simple architecture reduces computational overhead while still maintaining high predictive performance, which is essential for real-time applications in agriculture.

Overall, the proposed CNN-based system offers a well-rounded solution by combining high accuracy, fast processing speed, and computational efficiency. The model is designed to be scalable, allowing for future improvements and extensions without significant redesign. Its practical design makes it highly suitable for integration into mobile or web-based platforms, enabling farmers to easily use the system in their daily activities. By simply capturing an image of a plant leaf using a smartphone, users can obtain instant predictions about potential diseases. This capability supports early detection, timely intervention, and informed decision-making, ultimately reducing crop losses and improving agricultural productivity. The system not only enhances efficiency but also empowers farmers by providing them with accessible and reliable tools for crop monitoring.

While the proposed system demonstrates strong performance and practical usability, there are several opportunities for further improvement and enhancement to make it even more effective and versatile. One important area of improvement is the expansion of the dataset. Currently, the model is primarily trained on the PlantVillage dataset, which consists of images captured under controlled conditions with clean backgrounds. Although this helps achieve high accuracy, it does not fully represent real-world farming environments. In the future, incorporating images collected directly from fields, including variations in lighting, weather conditions, background clutter, and occlusions, will significantly improve the model's ability to generalize. This will make the system more robust and reliable when deployed in real agricultural settings.

Another potential enhancement is to increase the range of crops and diseases covered by the system. At present, the model is limited to a specific set of crops and disease categories. Expanding the dataset to include more crop varieties, additional disease types, and different stages of plant growth will make the system applicable to a broader audience. This will enable farmers from different regions and agricultural practices to benefit from the technology, thereby increasing its overall impact.

The integration of the system with smart farming technologies represents another promising direction for future development. By combining the model with tools such as drones, satellite imaging systems, and ground-based sensors, it would be possible to monitor large agricultural areas automatically. This would reduce the need for manual inspection and enable continuous surveillance of crop health across entire fields. Such integration can support precision agriculture practices, improve resource management, and enhance overall farm productivity.

In addition to disease detection, the system can be further enhanced by incorporating decision-support features. Currently, the model focuses on identifying diseases, but it can be extended to provide actionable recommendations, such as suggesting appropriate pesticides, organic treatments, or preventive measures. This would transform the system from a diagnostic tool into a comprehensive agricultural assistant, helping farmers not only identify problems but also take effective corrective actions.

Furthermore, advanced research techniques can be explored to improve the system's performance and adaptability. For instance, federated learning can be used to train models collaboratively across multiple devices while preserving data privacy. Multimodal learning approaches that combine image data with other sources such as soil conditions, weather data, and sensor readings can provide more accurate and context-aware predictions. Additionally, developing energy-efficient models that can run on low-power devices will make the system more sustainable and accessible, especially in rural areas with limited resources.

In conclusion, this study demonstrates the effectiveness of deep learning in addressing real-world agricultural challenges, particularly in the domain of plant disease detection. The proposed CNN-based system successfully combines accuracy, efficiency, and usability, making it a practical solution for modern farming. With further enhancements and integration of emerging technologies, the system has the potential to evolve into a comprehensive smart farming platform. Such advancements can significantly contribute to sustainable agriculture, improve crop management practices, and support farmers in achieving better productivity and economic outcomes.

REFERENCES

[1] T. B. Shahi, C.-Y. Xu, A. Neupane, and W. Guo, "Recent advances in crop disease detection using UAV and deep learning techniques," *Remote Sensing*, vol. 15, no. 9, p. 2450, 2023.

[2] A. Singla, A. Nehra, K. Joshi, A. Kumar, N. Tuteja, R. K. Varshney, S. S. Gill, and R. Gill, "Exploration of machine learning approaches for automated crop disease detection," *Current Plant Biology*, vol. 40, p. 100382, 2024.

[3] H. N. Ngugi, A. E. Ezugwu, A. A. Akinyelu, and L. Abualigah, "Revolutionizing crop disease detection with computational deep learning: A comprehensive review," *Environmental Monitoring and Assessment*, vol. 196, p. 302, 2024.

[4] Y. E. García-Vera, A. Poloche-Arango, C. A. Mendivelso-Fajardo, and F. J. Gutierrez-Bernal, "Hyperspectral image analysis and machine learning techniques for crop disease detection and identification: A review," *Sustainability*, vol. 16, no. 14, p. 6064, 2024.

[5] Y. Wang, U. S. R. Dhamodharan, N. Sarwar, F. A. Almalki, Q. H. Naith, R. Sathiyaraj, and D. Mohan, "A hybrid approach for rice crop disease detection in agricultural IoT systems," *Discover Sustainability*, vol. 5, p. 99, 2024.

[6] D. M. Kabala, A. Hafiane, L. Bobelin, and R. Canals, "Image-based crop disease detection with federated learning," *Scientific Reports*, vol. 13, p. 19220, 2023.

[7] J. Feng, W. E. Ong, W. C. Teh, and R. Zhang, "Enhanced crop disease detection with EfficientNet convolutional group-wise transformer," *IEEE Access*, vol. 12, pp. 44147–44162, 2024.

[8] H. Orchi, M. Sadik, M. Khaldoun, and E. Sabir, "Automation of crop disease detection through conventional machine learning and deep transfer learning approaches," *Agriculture*, vol. 13, no. 2, p. 352, 2023.

[9] Y. Wang, Y. Chen, and D. Wang, "Convolution network enlightened transformer for regional crop disease classification," *Electronics*, vol. 11, no. 19, p. 3174, 2022.

[10] R. Reedha, E. Dericquebourg, R. Canals, and A. Hafiane, "Transformer neural network for weed and crop classification of high-resolution UAV images," *Remote Sensing*, vol. 14, no. 3, p. 592, 2022.

[11] Y. Borhani, J. Khoramdel, and E. Najafi, "A deep learning-based approach for automated plant disease classification using vision transformer," *Scientific Reports*, vol. 12, p. 11554, 2022.

[12] X. Fu, Q. Ma, F. Yang, C. Zhang, X. Zhao, F. Chang, and L. Han, "Crop pest image recognition based on the improved ViT method," *Information Processing in Agriculture*, 2023.

[13] S. R. Dubey, S. K. Singh, and B. K. Singh, "CNN-based deep learning models for plant leaf disease classification: A review," *Computers and Electronics in Agriculture*, vol. 198, p. 107062, 2022.

[14] M. Kamilaris and F. X. Prenafeta-Boldú, "Deep learning in agriculture: A survey," *Computers and Electronics in Agriculture*, vol. 147, pp. 70–90, 2018.

[15] A. Sharma, A. Jain, P. Gupta, and V. Chowdary, "Machine learning applications for precision agriculture: A comprehensive review," *IEEE Access*, vol. 9, pp. 149443–149468, 2021.

[16] S. Sladojevic, M. Arsenovic, A. Anderla, D. Culibrk, and D. Stefanovic, "Deep neural networks based recognition of plant diseases by leaf image classification," *Computational Intelligence and Neuroscience*, vol. 2016, Article ID 3289801, 2016.

[17] A. Mohanty, D. Hughes, and M. Salathé, "Using deep learning for image-based plant disease detection," *Frontiers in Plant Science*, vol. 7, p. 1419, 2016.

[18] D. P. Hughes and M. Salathé, "An open access repository of images on plant health to enable the development of mobile disease diagnostics," *arXiv preprint arXiv:1511.08060*, 2015.

[19] K. P. Ferentinos, "Deep learning models for plant disease detection and diagnosis," *Computers and Electronics in Agriculture*, vol. 145, pp. 311–318, 2018.

[20] J. Amara, B. Bouaziz, and A. Algergawy, "A deep learning-based approach for banana leaf diseases classification," *Lecture Notes in Informatics*, 2017.

[21] P. Brahimi, A. Boukhalfa, and K. Moussaoui, "Deep learning for tomato diseases: Classification and symptoms visualization," *Applied Artificial Intelligence*, vol. 31, no. 4, pp. 299–315, 2017.

[22] S. Too, L. Yujian, S. Njuki, and L. Yingchun, "A comparative study of fine-tuning deep learning models for plant disease identification," *Computers and Electronics in Agriculture*, vol. 161, pp. 272–279, 2019.

[23] M. Rangarajan, R. Purushothaman, and A. Ramesh, "Tomato crop disease classification using pre-trained deep learning algorithm," *Procedia Computer Science*, vol. 133, pp. 1040–1047, 2018.

[24] S. Ramesh, D. Vydeki, and S. J. Raj, "Recognition and classification of paddy leaf diseases using optimized deep neural network," *Artificial Intelligence in Agriculture*, vol. 2, pp. 1–11, 2018.

[25] A. Kamilaris and F. X. Prenafeta-Boldú, "A review of the use of convolutional neural networks in agriculture," *Journal of Agricultural Science*, vol. 156, no. 3, pp. 312–322, 2018.

[26] Z. Zhang, H. Zhang, and X. Wang, "Plant disease recognition based on plant leaf image," *Journal*

of Ambient Intelligence and Humanized Computing, vol. 10, pp. 1–10, 2019.
[27] Y. LeCun, Y. Bengio, and G. Hinton, “Deep learning,” *Nature*, vol. 521, pp. 436–444, 2015.
[28] K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
[29] M. Tan and Q. Le, “EfficientNet: Rethinking model scaling for convolutional neural networks,” in *Proceedings of ICML*, 2019.
[30] A. Dosovitskiy et al., “An image is worth 16x16 words: Transformers for image recognition at scale,” in *International Conference on Learning Representations (ICLR)*, 2021.
[31] G. Huang, Z. Liu, L. Van Der Maaten, and K. Q. Weinberger, “Densely connected convolutional networks,” in *CVPR*, 2017.
[32] S. Albawi, T. A. Mohammed, and S. Al-Zawi, “Understanding of a convolutional neural network,” in *International Conference on Engineering and Technology*, 2017.
[33] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*, MIT Press, 2016.
[34] R. Girshick, “Fast R-CNN,” in *Proceedings of the IEEE International Conference on Computer Vision (ICCV)*, 2015.
[35] J. Redmon et al., “You Only Look Once: Unified, real-time object detection,” in *CVPR*, 2016.

APPENDIX 1

An Image-Based Study on Crop Disease Detection

Mohammad Areeb[1], Ayush Tripathi[1], Aman Bohra[1], Yash Vinayak Singh[1], Rahul Kumar Sharma[1], Devendra Gautam[2]

[1]Department of Computer Science and Engineering, Krishna Institute of Engineering & Technology (KIET), Ghaziabad, Delhi-NCR, Uttar Pradesh, India

mohammad.2226cse1105@kiet.edu,ayush.2226cse1146@kiet.edu,aman.2226cse1146@kiet.edu, yash.2226cse1165@kiet.edu, rahul.sharma@kiet.edu

[2]JIMS EMTC, Greater Noida, India

devendra.gautamcs16@gmail.com

***Abstract*— Crop diseases is a major problem for farmers all around the world as it results in reduced yield and losses in billions of dollars. To contribute in solving this problem we developed an automated crop disease detection system which can detect diseases by analyzing a photograph of a leaf. We have implemented a deep learning-based system to classify crop diseases. We have trained a Convolutional Neural Network (CNN) based model on PlantVillage Dataset containing a total of 61,486 images across 38 classes (each class corresponds to a type of disease with a class for healthy leaves as well). We have pre-processed and augmented the data before training it on the model. This is done to improve the accuracy of our model in real world scenarios. Steps like resizing the images to a fixed size, cropping the images, rotating the images, varying brightness etc. were implemented as part of preprocessing and augmentation steps. The model is then used to visually analyse the diseased portions on the leaf and display the results quickly on a user interface. Model has a classification accuracy of 98.7% for images under control conditions as of initial testing (accuracy may vary with other factors like environment, background, lighting, quality of image etc.). Currently, the system is built using the CNN model. The plan is to build and train more complex and efficient models in the future using models like EfficientNet, ResNet, Vision Transformers and more. We also plan to use the trained model for a mobile application in the future to identify plant diseases in the field. The proposed system is cost-effective, fast, easy-to-use, and highly scalable which can potentially help in reducing crop losses and aiding farmers.**

Index Terms—**Crop Disease Detection, Convolutional Neural Network, PlantVillage, Deep Learning, Image Classification, Agriculture Technology**

I. Introduction

Agriculture plays a critical role in ensuring food security and driving economies across the world. Despite advancements in farming practices, plants are susceptible to a wide array of diseases. These diseases can reduce the yield of crops, degrade the quality of produce, and result in significant economic losses for farmers. According to agriculture reports, plant diseases account for 20% to 40% of crop losses every year. This alarming loss in productivity emphasizes the need for early detection and intervention. Traditionally, visual inspection of crops was the go-to method for identifying plant diseases. However, this manual approach can be time-consuming, prone to inaccuracies, and less effective in remote or rural areas with limited expertise. The consequence of late disease detection or misdiagnosis can be grave for the farmers, as it might lead to overuse of pesticides, crop wastage, and food shortage. The ad- vent of Artificial Intelligence (AI) and Computer Vision (CV) has opened new avenues for accurate and efficient plant disease detection. Deep learning, particularly Convolutional Neural Networks (CNNs), has emerged as a promising approach as it can automatically learn features from images to identify diseases. In this study, we have developed a deep learning- based system for the detection of 38 types of plant diseases and healthy leaves from the given input image. The system has been trained on a diverse PlantVillage dataset that comprises more than 61,000 images. With robust image preprocessing and CNN architecture, the system is fast, accurate, and can be deployed for general use. We hope to make advanced AI tools more accessible to the farmers and help them with the early identification of crop diseases. The main contributions of this work are - 1. Development of a lightweight CNN model for crop disease detection. 2. Extensive evaluation on the PlantVillage dataset containing 61,486 images.3. Comparison with modern architectures including MobileNetV2, ResNet50, and EfficientNet.

II. LITERATURE REVIEW

Efforts to find a suitable method for disease identification have been done in a rapidly expanding research space that investigates various techniques such as machine learning (ML), deep learning (DL), hyperspectral, federated learning (FL), and supervised and unsupervised learning schemes. While some works have evaluated various types of these methods to achieve high accuracy, others have also been more concerned with real-world applicability by tuning and validating their models to perform optimally in non-laboratory situations. An example of the former includes an effort by [12] in 2023, which conducts an analysis of ML methods such as support vector machine (SVM), k-nearest neighbors (KNN) with transfer learning models (VGG16, InceptionV3, and ResNet50) to determine which optimizer and Convolutional Neural Network (CNN) model configuration is best-suited to a certain dataset. The study successfully finds an accuracy of roughly 98.01% for InceptionV3 (Trained with Adaptive Moment Estimation (Adam) optimizer) using the PlantVillage dataset and experiments from multiple sources to further show the capabilities of transfer learning as a potential means for disease identification in multiclass scenarios. On the other hand, some others such as a study by [11] have further evaluated the potential of FL with CNNs and Vision Transformers as a possible solution to data privacy concerns and the limitations of centralized learning. The researchers of [10] successfully demonstrate a proof-of- concept by using a dynamic agri-ecosystem environment with a large distribution of agents and content providers. Their findings show that the number of agents or contributors to data, number of communication rounds, and other variables can impact the resulting model's performance.

A recent review by [3] demonstrated the progress in the DL and ML methods being used in the detection and recognition of plant diseases, covering CNNs, RNNs, SVMs, and various combinations of supervised, semi-supervised, and self-supervised learning (SSL) models. In addition to providing a comparison between the relative benefits and shortcomings of various methods and datasets, the review also discussed the limitations in generalizability and cross-compatibility between models in a similar manner to some others such as [8]. Other review papers such as one in [7] have also discussed an alternative means of identifying plant diseases using hyperspectral imaging in combination with AI methods. While its ability to perform spectral analysis of crops at such a fine granularity shows promise, there are existing issues with cost, large storage requirements, and the need for specialized and bulky hard- ware. More crop-specific works such as in [5] show that customizing the type of AI method to suit a problem is a possible future direction, proposing an architecture for smart farming disease identification in rice that combines DC-GAN and MDFCResNet.

These represent novel methodologies to enhance information extraction from images. The group-wise transformer-augmented EfficientNet work resulted in an accuracy of 97.6% by fusing features at various levels. Transfer learning-based solutions such as LeafDNet with Xception architecture achieved high accuracies of approximately 98% in rose leaf disease classification. YOLOv8 based initiatives are being developed for real-time disease detection, aimed at empowering farmers on-the-go. In comparison to the diverse directions taken in various studies, the CNN-based approach proposed in this research strikes a middle ground. It achieves competitive accuracy of 98.7% while remaining computationally efficient. It doesn't depend on costly hyperspectral imaging, intricate hybrid architectures, or federated networks. It features user-friendly and easy preprocessing steps and contemplates a mobile app for real-world applications, particularly in rural areas. Consequently, the proposed system presents itself as a promising candidate for affordable, practical, and effective broader deployment.

III. PROPOSED SYSTEM

In order to address the issues of manual verification and classical ML approach. The given project proposes a deep learning model for automatic detection of crop diseases. This system(fig. 1) can identify the disease on a plant rapidly and precisely by taking the image of its leaf as an input. The model is developed using a customized CNN which is trained on a PlantVillage dataset with 38 different diseases spread across multiple crops. Rather than extracting features that humans can understand, it can learn key features on its own such as leaf spots, blight, mosaic, mildew, color variations, etc.

A careful preprocessing and image augmentation (resizing, normalization, rotation, flipping, and color change) process ensures the system's robustness under varying light, orientation, and background conditions. The model returns the predicted disease and a confidence value for each input, which allows for high speed and reliable operation. The system is available both on the web and for mobile devices and allows farmers to snap photos of leaves for near- instantaneous feedback. The model is most useful for farmers in remote or resource-poor areas. There is also a visualization tool to see confidence levels, historical predictions, and disease trends which can be used for both day-to-day farming and large-scale agricultural monitoring.

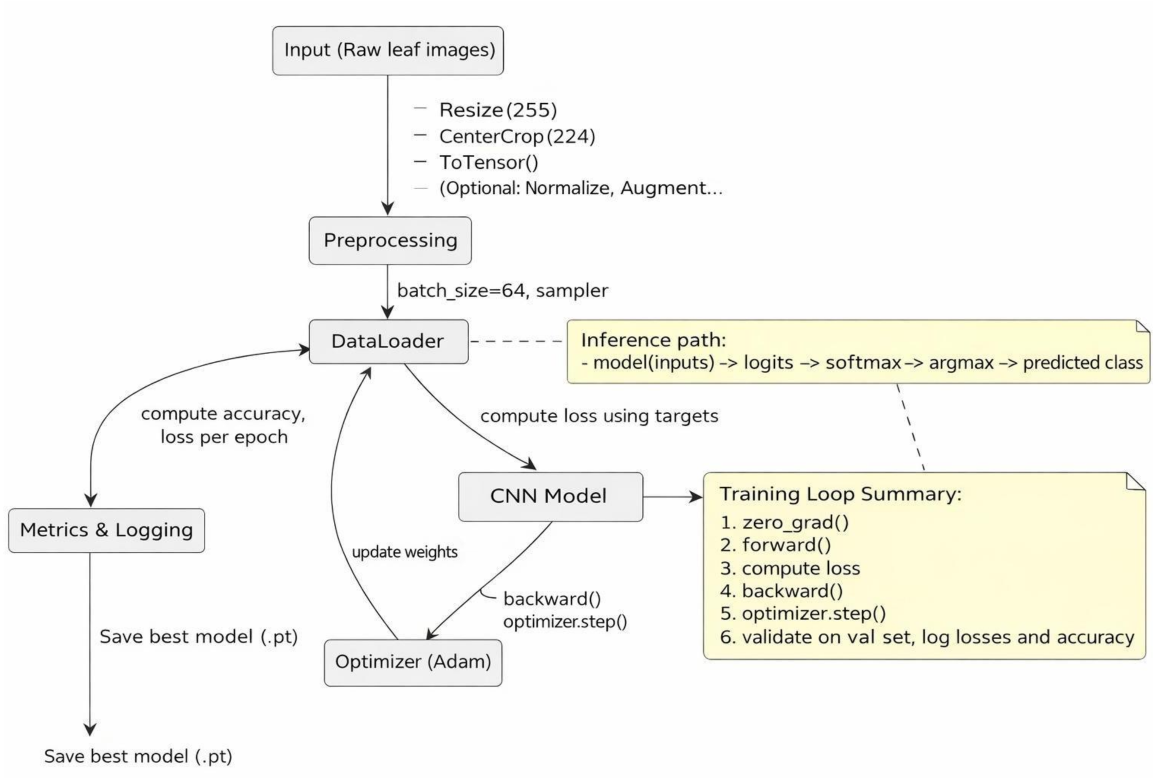

Fig.1 Proposed System

IV. METHODOLOGY

The proposed **Crop Disease Detection System** follows a modular deep learning pipeline integrating image preprocessing, CNN-based feature extraction, and real-time deployment for mobile environments. The architecture is designed to be lightweight, scalable, and accurate for multi-crop disease detection.

A. Requirement Analysis:

Traditional crop disease detection methods rely heavily on agricultural experts and manual inspection, resulting in high misdiagnosis rates and limited scalability. Conventional machine learning and hyperspectral approaches also suffer from poor generalization, handcrafted feature dependency, high computational requirements, and limited usability in mobile environments.

To overcome these limitations, a lightweight Convolutional Neural Network (CNN) based system was developed to enable fast and accurate disease identification directly on smartphones, allowing farmers to diagnose crop diseases in real-time under field conditions.

B. System Design:

The system architecture consists of three major modules:
1) Image Preprocessing Module:
Input leaf images are cleaned and standardized before being fed to the neural network.
2) Disease Classification Module (CNN):
A custom CNN architecture performs automatic feature extraction and multi-class classification across 38 plant disease categories.
(i) Convolution Operation: Extracts important spatial features such as edges, textures, spots, and colour variations from the input image using filters.
(ii) ReLU Activation: Introduces non-linearity by converting negative values to zero while retaining positive values.
(iii) Batch Normalization: Normalizes layer outputs during training to improve training stability and accelerate model convergence.
(iv) Max Pooling: Reduces the spatial dimensions of feature maps while preserving the most important features.
(v) Fully Connected Layer: Combines the extracted features from previous layers to perform the final classification.

(vi) SoftMax for 38-Class Output: Converts the network outputs into probability values for each of the 38 possible disease classes.
(vii) Cross-Entropy Loss: Measures the difference between predicted probabilities and the actual class labels, guiding the model during training.
3) Deployment and User Interface:
The trained model is deployed as a real-time diagnostic tool accessible via smartphones. Users capture a crop leaf image, which is analyzed locally or via cloud inference. The system returns the predicted disease and confidence score while maintaining historical predictions and visualization dashboards to track disease trends.

C. Datasets:

Two datasets were used.
PlantVillage Dataset

- 61,486 RGB leaf images
- 38 disease classes
- Split: 70% training, 15% validation, 15% testing

Field Dataset (Future Extension)

- Real farm images.
- Contains shadows, noise, and background clutter.
- Used for evaluating model generalization.

D. Feature Extraction:

CNN layers automatically extract hierarchical visual features:

- Low-level: edges, color gradients.
- Mid-level: lesions, spots.
- High-level: complex disease patterns.

These features enable detection of fungal infections, discoloration, pest damage, and viral symptoms.

E. Model Implementation:

1) Custom CNN Architecture:

- 3 × (Conv → ReLU → BatchNorm) blocks.
- Max pooling after each block.
- Dropout regularization.
- Flatten + Dense layer.
- SoftMax output (38 classes).

2) Transfer Learning Models:
To evaluate lightweight architectures suitable for mobile deployment, the following models were also tested:

- MobileNetV2.
- ResNet50.
- EfficientNet-B0.

3) Training Configuration:
Optimizer: Adam, Learning rate: 0.001, Batch size: 32, Epochs: 50, Loss function: Cross-Entropy Loss.

F. Pseudocode

1. Initialize system: Load trained CNN model and disease label set.
2. Input image: Capture or upload crop leaf image.
3. Preprocess image: Resize to 256×256, normalize, and apply augmentation (training phase).
4. Feature extraction: Apply convolution, ReLU, batch normalization, and max pooling layers.
5. Classification: Flatten feature maps and compute class probabilities using SoftMax.
6. Prediction: Select disease class with highest probability.
7. Output: Display predicted disease and confidence score through the mobile/web interface.

V. RESULT AND DISCUSSION

This part demonstrates the experimental results of the proposed CNN-based method with the analysis of training process, performance on validation set(fig. 2) and comparisons with other architectures(Table 1). The findings show the model's stability, its good generalization ability and its potential to be utilized in actual agricultural environment. The accuracy curves for both training and validation showed smooth and stable convergence. The model reached 99.4% accuracy on the training set and 98.6% on the validation set. The loss values kept going down with each training round, and the small difference between training and validation loss shows that the model didn't overfit much. Overall, these results show that the custom CNN successfully picked up important patterns related to the diseases from the varied PlantVillage dataset.

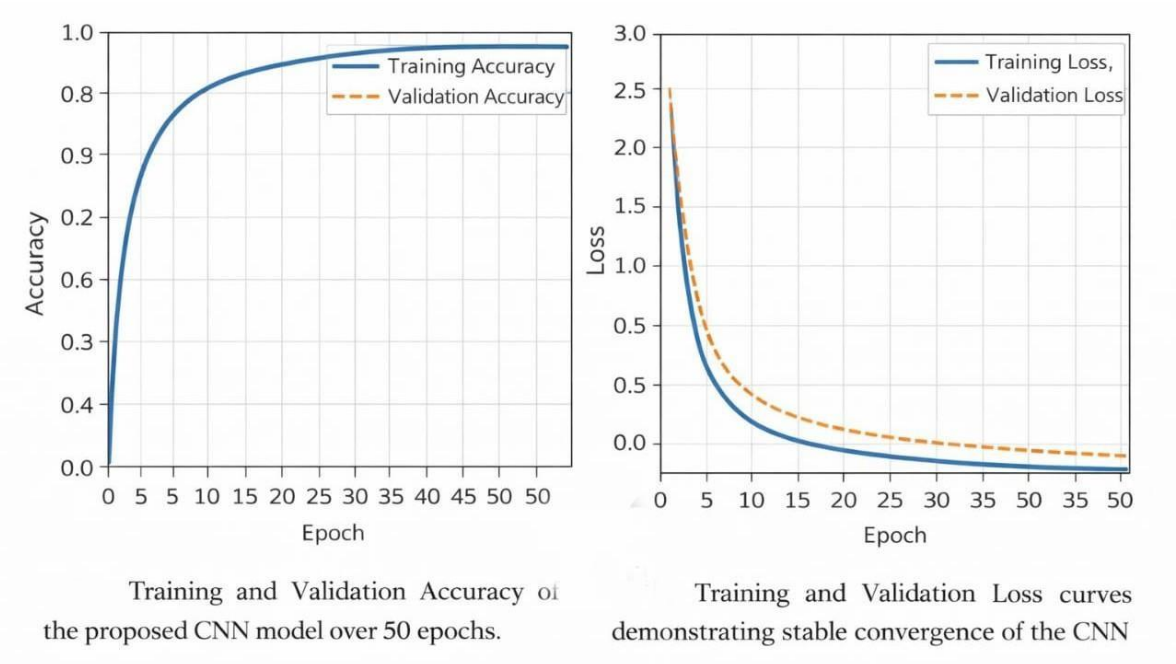

Fig. 2 Training and Validation Accuracy & Loss Graph

Architecture	Test Accuracy	Remarks
Proposed CNN	98.7%	Achieved the highest accuracy among all evaluated models
MobileNetV2	96.3%	Optimized for low-resource and mobile-based applications
ResNet50	97.5%	Effective feature learning with minor overfitting tendencies
EfficientNet-B0	97.9%	Good trade-off between accuracy and computational cost

Table 1 Performance Evaluation of Different Deep Learning Architecture

VI. CONCLUSION

This study investigated crop disease detection using deep learning techniques. A custom Convolutional Neural Network (CNN) was trained using the PlantVillage dataset and its performance was compared with MobileNetV2, ResNet50, and EfficientNet-B0 architectures. The proposed CNN achieved the highest test accuracy of **98.7%**, demonstrating strong generalization capability and consistent performance across different disease classes.

MobileNetV2 achieved slightly lower accuracy but remains a suitable lightweight architecture for mobile and resource-constrained environments. ResNet50 and EfficientNet-B0 also demonstrated strong feature extraction capabilities; however, both models showed minor overfitting when evaluated on limited real-world field images.

Overall, the proposed CNN model provides **fast, accurate, and scalable crop disease detection**, making it a practical solution for real-world agricultural applications. The system can serve as a key component in **mobile-based farmer assistance platforms**, enabling timely disease diagnosis and improving crop health management.

VII. FUTURE SCOPE

Although the proposed crop disease detection system demonstrates strong performance, several improvements can further enhance its applicability and robustness:

• **Dataset Expansion:** Increase the dataset with additional real-world field images captured under diverse lighting conditions, weather variations, and complex backgrounds to improve model generalization.

• **Broader Crop Coverage:** Extend the dataset to include more crop varieties, additional disease categories, and different plant growth stages, enabling wider agricultural applicability.

• **Integration with IoT and Smart Farming:** Integrate the system with drones, satellite imagery, and ground-based sensors to enable automated large-scale crop health monitoring.

• **Decision-Support Features:** Enhance the system with actionable recommendations including **pesticide suggestions, biological treatments, and preventive measures** to assist farmers in effective disease management.

• **Advanced Research Directions:** Explore advanced approaches such as **federated learning for privacy-preserving training, multimodal analysis combining image and sensor data, and energy-efficient inference methods** for sustainable precision agriculture.

VIII. REFERENCES

[1] T. B. Shahi, C.-Y. Xu, A. Neupane, and W. Guo, "Recent advances in crop disease detection using UAV and deep learning techniques," *Remote Sensing*, vol. 15, no. 9, p. 2450, 2023.

[2] A. Singla, A. Nehra, K. Joshi, A. Kumar, N. Tuteja, R. K. Varshney, S. S. Gill, and R. Gill, "Exploration of machine learning approaches for automated crop disease detection," *Current Plant Biology*, vol. 40, p. 100382, 2024.

[3] H. N. Ngugi, A. E. Ezugwu, A. A. Akinyelu, and L. Abualigah, "Revolutionizing crop disease detection with computational deep learning: A comprehensive review," *Environmental Monitoring and Assessment*, vol. 196, p. 302, 2024.

[4] Y. E. García-Vera, A. Poloche-Arango, C. A. Mendivelso-Fajardo, and F. J. Gutierrez-Bernal, "Hyperspectral image analysis and machine learning techniques for crop disease detection and identification: A review," *Sustainability*, vol. 16, no. 14, p. 6064, 2024.

[5] Y. Wang, U. S. R. Dhamodharan, N. Sarwar, F. A. Almalki, Q. H. Naith, R. Sathiyaraj, and D. Mohan, "A hybrid approach for rice crop disease detection in agricultural IoT systems," *Discover Sustainability*, vol. 5, p. 99, 2024.

[6] D. M. Kabala, A. Hafiane, L. Bobelin, and R. Canals, "Image-based crop disease detection with federated learning," *Scientific Reports*, vol. 13, p. 19220, 2023.

[7] J. Feng, W. E. Ong, W. C. Teh, and R. Zhang, "Enhanced crop disease detection with EfficientNet convolutional group-wise transformer," *IEEE Access*, vol. 12, pp. 44147–44162, 2024.

[8] H. Orchi, M. Sadik, M. Khaldoun, and E. Sabir, "Automation of crop disease detection through conventional machine learning and deep transfer learning approaches," *Agriculture*, vol. 13, no. 2, p. 352, 2023.

[9] Y. Wang, Y. Chen, and D. Wang, "Convolution network enlightened transformer for regional crop disease classification," *Electronics*, vol. 11, no. 19, p. 3174, 2022.

[10] R. Reedha, E. Dericquebourg, R. Canals, and A. Hafiane, "Transformer neural network for weed and crop classification of high-resolution UAV images," *Remote Sensing*, vol. 14, no. 3, p. 592, 2022.

[11] Y. Borhani, J. Khoramdel, and E. Najafi, "A deep learning-based approach for automated plant disease classification using vision transformer," *Scientific Reports*, vol. 12, p. 11554, 2022.

[12] X. Fu, Q. Ma, F. Yang, C. Zhang, X. Zhao, F. Chang, and L. Han, "Crop pest image recognition based on the improved ViT method," *Information Processing in Agriculture*, 2023.

PROOF OF SUBMISSION/ACCEPTANCE

INTERNATIONAL CONFERENCE ON SUSTAINABLE AI FOR CYBERSECURITY

CERTIFICATE OF PRESENTATION

This certificate is proudly presented to

Aman Singh Bohra

For presenting the paper titled

An Image - Based Study on Crop Disease Detection

at the International Conference on Sustainable AI for Cybersecurity organized by the Department of Computer Science and Engineering (CSE) at JIMS Greater Noida in collaboration with Hinweis Research, held during 8th - 9th April 2026.

Prof. (Dr.) Latha Banda
HOD CSE, JIMSEMTEC

Prof. (Dr.) Sachin Yadav
Director, JIMSEMTC

Dr. Amit Gupta
Chairman, JIMS Group of Institutions